

# CE 440 Introduction to Operating System

## Lecture 5: Scheduling Fall 2025

Prof. Yigong Hu

BOSTON  
UNIVERSITY

Slides courtesy of Manuel Egele, Ryan Huang and Baris Kasikci

# Administrivia

## Lab 0

- Due this Friday
- Done individually (cannot share with or copy from your to-be-teammates)

## Find your project group member soon

- So you can get started with Lab 1 without delay
- Fill out Google form of group info ( will upload on Piazza)
  - [https://docs.google.com/forms/d/e/1FAIpQLScqr0QdmoruMu\\_w7-FizeQ9OYaijg9-d9Y58zOV28wivnYp5A/viewform?usp=dialog](https://docs.google.com/forms/d/e/1FAIpQLScqr0QdmoruMu_w7-FizeQ9OYaijg9-d9Y58zOV28wivnYp5A/viewform?usp=dialog)

# Recap: Processes, Threads

**Process is the OS abstraction for execution**

- own view of machine

**Process components**

- address space, program counter, registers, open files, etc.
- kernel data structure: **Process Control Block (PCB)**


**Process vs. thread**

**Process/thread states and APIs**


- state graph and queues
- process creation, deletion, waiting

**Multiple processes/threads**

- overlapping I/O and CPU activities
- context switch



# Scheduling Overview



**The scheduling problem:**

- Have  $K$  jobs ready to run
- Have  $N \geq 1$  CPUs

**Policy:** which jobs should we assign to which CPU(s), for how long?

- we'll refer to schedulable entities as **jobs** – could be processes, threads, people, etc.

**Mechanism:** context switch, process state queues

# Scheduling Goals

## Goal 1: guarantee “good service”

- To decide what job to run next and for how long
- Good service could be one of many different criteria
  - Fairness – giving each process a fair share of the CPU
  - Throughput – maximize jobs per second
  - Response time - respond to requests quickly

## Known as short-term scheduling decision

- Happens relatively **frequently**
- Want to minimize the overhead of scheduling
  - Fast context switches, fast queue manipulation

# Scheduling Goals

## Goal 2: loaded jobs into memory

- To determine the multiprogramming level: how many jobs to run simultaneously
- Moving jobs to/from memory is often called swapping

## Known as long-term scheduling decision

- Happens relatively infrequently
- Significant overhead in swapping a process out to disk

Virtual Memory Lecture (Lecture 10-13)

# What Is “Good Service”?

**How do we measure the effectiveness of a scheduling algorithm?**

## Batch systems strive for

- Throughput – # of processes that complete per unit time
  - $\# \text{jobs}/\text{time}$
  - Higher is better
- Turnaround time – time for each process to complete
  - $T_{\text{finish}} - T_{\text{start}}$
  - Lower is better
- $CPU \text{ utilization}$  –  $\%CPU$  fraction of time CPU doing productive work

# What Is “Good Service”?

## Interactive systems strive to

- minimize response time for interactive jobs (PC)
  - $T_{response} - T_{request}$ : time between *waiting* → *ready* transition and *ready* → *running*
  - Lower is better
- Proportionality – meet users’ expectations
  - Service-level objective(SLO)
- Utilization and throughput are often traded off for better response time

## Real-time systems

- Meeting deadlines: avoid losing data
- Predictability: avoid quality degradation in multimedia systems

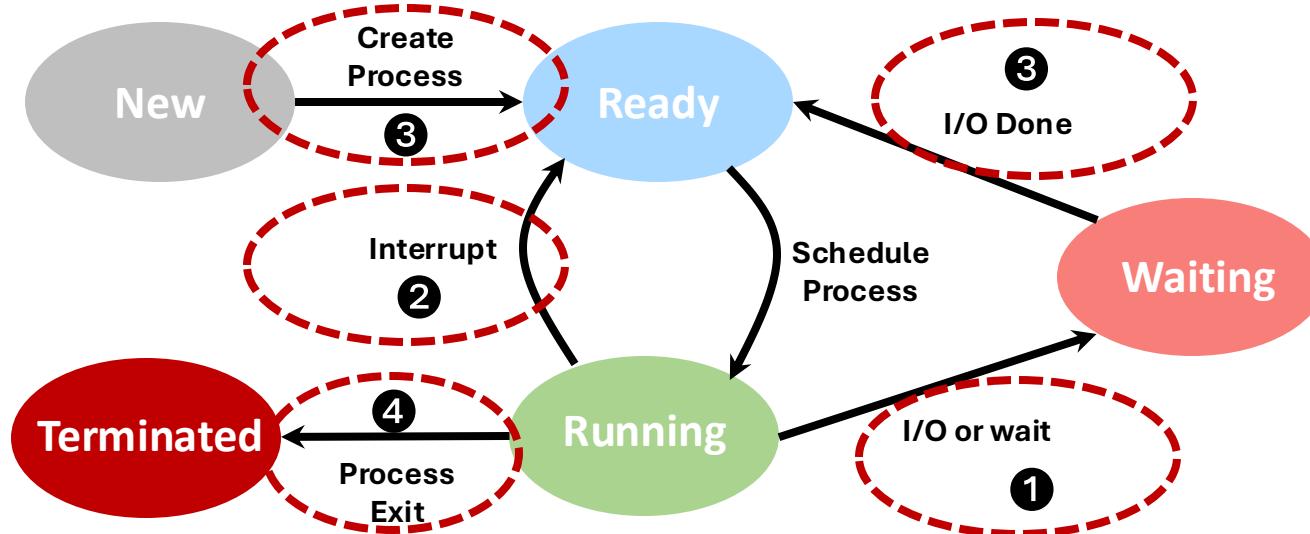
# Tradeoffs

**Improving on one metric can hurt another**

**For example:**

- We want to improve throughput, so we decide to only schedule short jobs
- But now longer jobs never get run, so their turnaround time is effectively infinite

# When Do We Schedule CPU?



**Scheduling decisions may take place when a process:**

- ① Switches from running to waiting state
- ② Switches from running to ready state
- ③ Switches from new/waiting to ready
- ④ Exits

**Non-preemptive schedules use ① & ④ only**

**Preemptive schedulers run at all four points**

# Scheduling Overviews

- Textbook scheduling
- Priority scheduling
- Advanced scheduling topics (not covered)

# FCFS Scheduling

“First-come first-served” (FCFS): Run jobs in order that they arrive

## Examples:

- Say P1 needs 24 sec, while P2 and P3 need 3.
- Say P2, P3 arrived immediately after P1



Throughput: 3 jobs / 30 sec = 0.1 jobs/sec

Turnaround Time: P1 : 24, P2 : 27, P3 : 30

- Average TT:  $(24 + 27 + 30) / 3 = 27$

Waiting Time: P1 : 0, P2 : 24, P3 : 27

- Average WT:  $(0 + 24 + 27) / 3 = 17$

Can we do better with FCFS?

# FCFS Scheduling Continued

Suppose we scheduled P2, P3, then P1



Throughput: 3 jobs / 30 sec = 0.1 jobs/sec

Turnaround Time: P1 : 30, P2 : 3, P3 : 6

- Average TT:  $(30 + 3 + 6) / 3 = 13$

Observations: scheduling algorithm can reduce TT

- Minimizing waiting time can improve RT and TT

Can a scheduling algorithm improve throughput?

- Yes, if jobs require both computation and I/O

# Scheduling Jobs with Computation & I/O

**CPU is one of several devices needed by users' jobs**

- CPU runs compute jobs, Disk drive runs disk jobs, etc.
- With network, part of job may run on remote CPU


**Scheduling 1-CPU system with n I/O devices like scheduling asymmetric (n + 1)-CPU multiprocessor**

- Result:  $(n + 1)$ -fold throughput gain!

# Scheduling Jobs with Computation & I/O(2)

Example: **disk-bound** grep + **CPU-bound** matrix\_multiply

- Overlap them just right, throughput will be almost doubled




# FCFS Limitations

**FCFS algorithm is non-preemptive in nature**

- Once CPU time has been allocated to a process, other processes can get CPU time only after the current process has finished or gets blocked.

This property of FCFS scheduling is called ***Convoy Effect***



# Shortest Job First (SJF)

## Shortest Job First (SJF)

- Choose the job with the smallest expected CPU burst
- Person with smallest # of items in shopping cart checks out first

## Examples:

- Say P1 needs 8 sec, P2 4 sec and P3 2 sec.



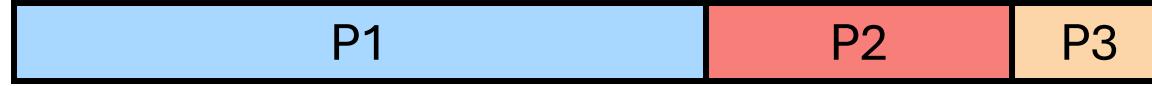
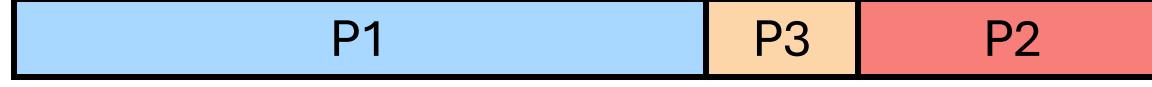
Average Waiting Time:  $(0 + 2 + 6) / 3 = 2.67$

# SJF Has Optimal Average Waiting Time

SJF has *provably* optimal minimum average waiting time (AWT)

## Previous Examples:

- P1 needs 8 sec, P2 4 sec and P3 2 sec.

Schedule 1		AWT: $(0 + 8 + 12) / 3 = 6.67$
Schedule 2		AWT = $(0+8+10)/3 = 6$
Schedule 3		AWT = $(0+4+12)/3 = 5.33$
Schedule 4		AWT = $(0+4+6)/3 = 3.33$
Schedule 5		AWT = $(0+2+10)/3 = 4$
SJF		AWT = $(0+2+6)/3 = 2.67$


Problem: what if new jobs arrive?

# Counterexample

The optimality proof only applies when all jobs are available at time 0

Suppose we have instead:

- At time 0, P1 needs 4 sec and P2 needs 5 sec.
- At time 2 seconds, processes P3, P4, and P5 arrive, each requiring 1 second of CPU time.

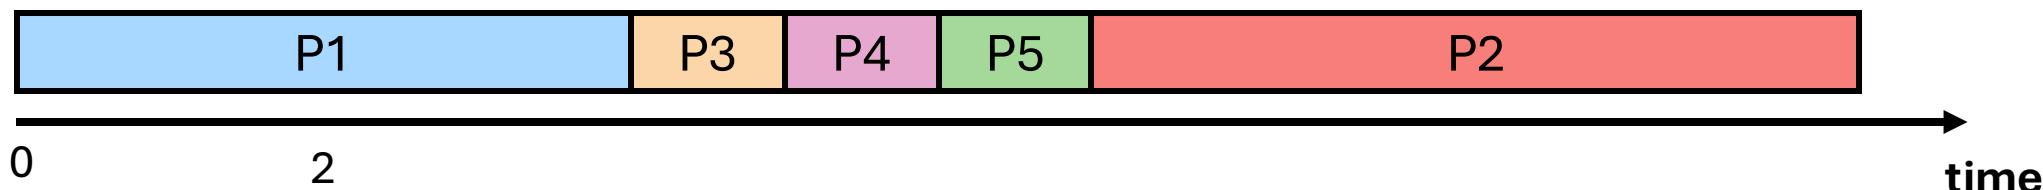


# Counterexample

The optimality proof only applies when all jobs are available at time 0

Suppose we have instead:

- At time 0, P1 needs 4 sec and P2 needs 8 sec.
- At time 2 seconds, processes P3, P4, and P5 arrive, each requiring 1 second of CPU time.



What is the AWT?

# Shortest Remaining Time Next

**SRTF chooses the process whose remaining run time is the shortest**


- When a new job arrives, its remaining run time is compared to the one of the currently running process
- If current process has more remaining time than the run time of new process, the current process is **preempted** and the new one is run

# Examples with Preemptive

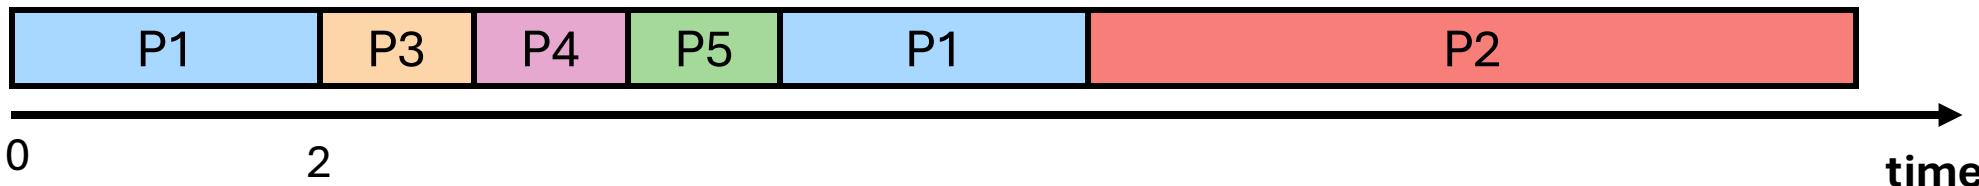
Process	Arrive Time	Burst Time
P1	0	4
P2	0	5
P3	2	1
P4	2	1
P5	2	1

What is the AWT?

Non-preemptive SJF:



Preemptive SRJF:



# SJF Limitations

**This algorithm also assumes that running time for all the processes to be run is known in advance**

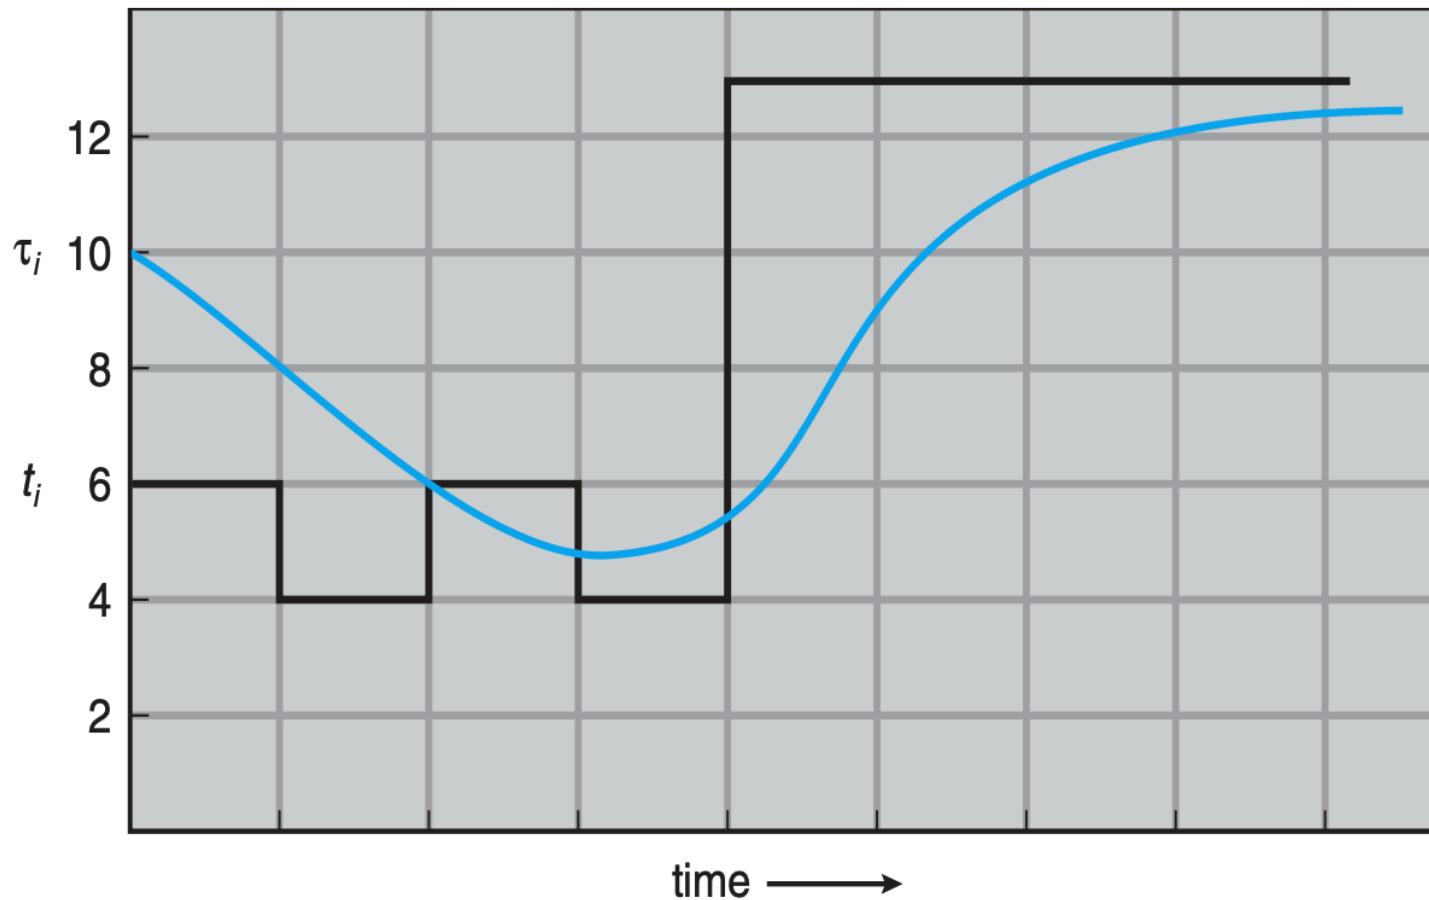
- Impossible to know size of CPU burst ahead of time

**Can potentially lead to unfairness or starvation**

**How can you make a reasonable guess?**

- Estimate CPU burst length based on past
- E.g., exponentially weighted average
  - $t_n$  actual length of process's  $n^{th}$  CPU burst
  - $\tau_{n+1}$  estimated length of proc's  $(n + 1)^{st}$  CPU burst
  - Choose parameter  $\alpha$  where  $0 < \alpha \leq 1$ , e.g.,  $\alpha = 0.5$
  - Let  $\tau_{n+1} = t_n + (1 - \alpha) \tau_n$

# Exp. Weighted Average Example



CPU burst ( $t_i$ )

6 4 6 4 13 13 13 ...

"guess" ( $\tau_i$ )

10 8 6 6 5 9 11 12 ...

# Round Robin (RR)

**Now, since we have preemptive scheduling:**

- Each process gets a small unit of CPU time (time quantum), usually 10-100 milliseconds
- Run first process until its quantum is used up
- Move that process to the end and run the next process
- Simple, fair
  - No process waits forever

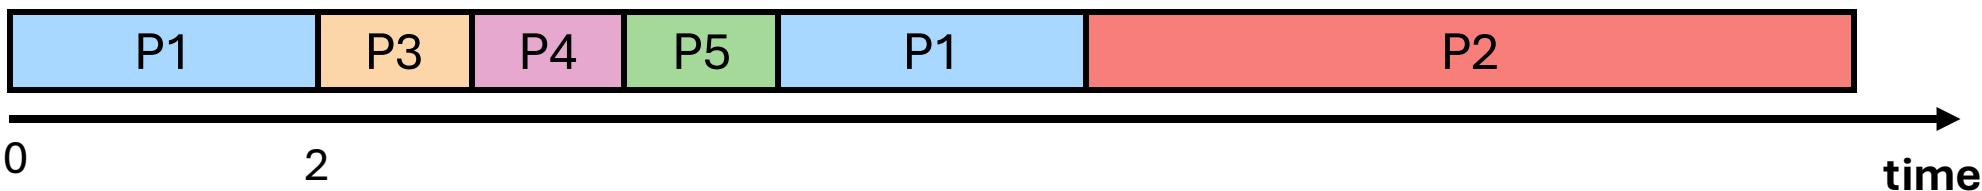
**Solution to fairness and starvation**

- Each job is given a time slice called a **quantum**
- Preempt job after duration of quantum
- When preempted, move to back of FIFO queue

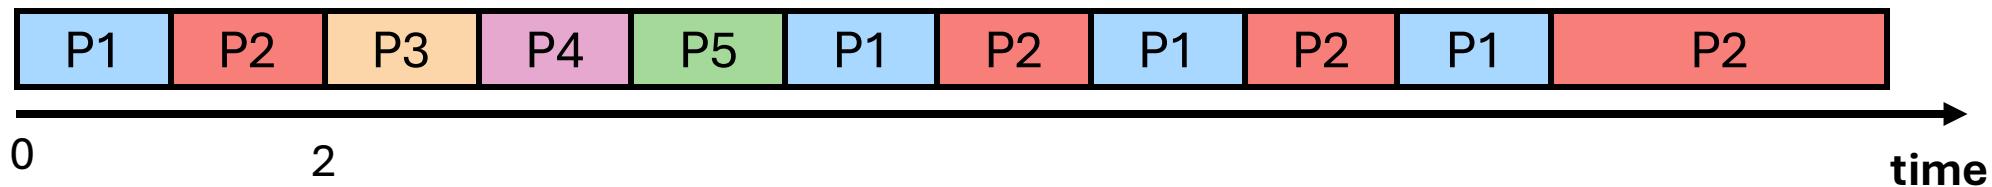
# Examples with Round Robin

Process	Arrive Time	Burst Time
P1	0	4
P2	0	5
P3	2	1
P4	2	1
P5	2	1

Preemptive SRJF:



Round Robin with quantum as 1 second



# **Advantage of Round Robin**

## **Solution to fairness and starvation**

- Each job is given a time slice called a **quantum**
- Preempt job after duration of quantum
- When preempted, move to back of FIFO queue

## **Advantages:**

- Fair allocation of CPU across jobs
- Low average waiting time when job lengths vary
- Good for responsiveness if small number of jobs

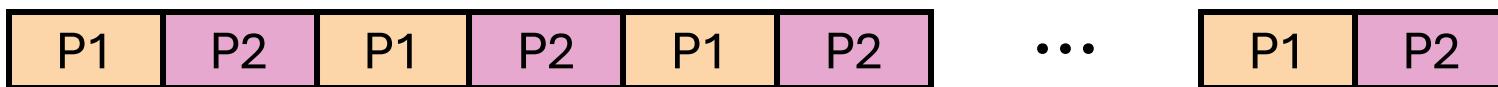
## **Disadvantages?**

# Disadvantages of Round Robin

**Context switches are frequent and need to be very fast**

**Varying sized jobs are good ...what about same-sized jobs?**

**Assume 2 jobs of time=100 each:**



**Even if context switches were free...**

- What would average turnaround time be with RR?
- Even worse than FCFS

# Round Robin Discussion

## How to pick quantum?

- What if too big?
  - Response time can be very bad
- What if time slice too small?
  - A notable percentage of the CPU time is spent in switching contexts

## Actual choices of time slice:

- Initially, UNIX time slice one second:
  - Worked ok when UNIX was used by one or two people.
  - What if three compilations going on? 3 seconds to echo each keystroke!
- Need to balance short-job performance and long-job throughput
  - Typical time slice today is between **10ms – 100ms**

# Scheduling Overviews

- Textbook scheduling
- Priority scheduling
- Advanced scheduling topics (not covered)

# Priority Scheduling

## Priority Scheduling

- Associate a numeric priority with each process
  - E.g., smaller number means higher priority (Unix/BSD)
  - Or smaller number means lower priority (Pintos)
- Give CPU to the process with highest priority
  - Airline check-in for first class passengers
  - Can be done preemptively or non-preemptively
- Can implement SJF, priority =  $1/(\text{expected CPU burst})$

**Problem:** starvation – low priority jobs can wait indefinitely


## Solution? “Age” processes

- Increase priority as a function of waiting time
- Decrease priority as a function of CPU consumption


# Examples with Priority Scheduling

Process	Arrive Time	Burst Time	Priority
P1	0	5	2
P2	3	1	1
P3	4	3	4
P4	8	7	0
P5	12	2	3

**Non-preemptive priority scheduling:**



**Preemptive priority scheduling**




# Priority Inversion (1)

## Caveat using Priority Scheduling w/ Synch Primitives

- Priority scheduling rule
  - 1) Always pick highest-priority thread
  - 2) ...unless a lower-priority thread is holding a resource the highest-priority thread wants to get
- Potential *Priority Inversion* Problem

Two tasks: **H** at high priority, **L** at low priority




# Priority Inversion (2)

**Two tasks: *H* at high priority, *L* at low priority**

- What if we have a task *M* enters system at medium priority, preempts *L*
- *L* unable to release R in time, *H* unable to run, despite having higher priority than *M*

**Not just a hypothetical issue, it happened in real-world software!**

- The root cause for a famous Mars Pathfinder failure in 1997
- Low-priority data gathering task and a medium-priority communications task prevented the critical bus management task from running



# Solution: Priority Donation

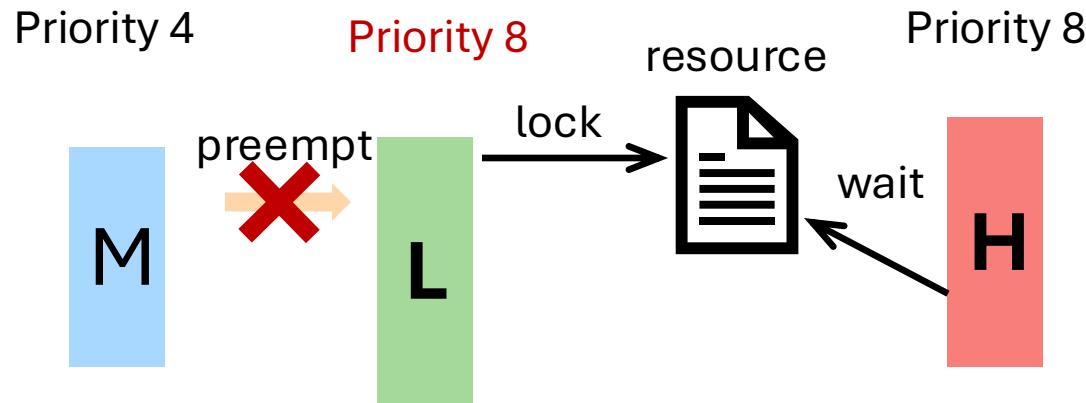
## “Donate” our priority if we get blocked

- Whenever a high-priority task has to wait for some shared resource that currently held by an executing low priority task,
- the low-priority task is *temporarily* assigned the priority of the highest waiting priority task for the duration of its use of the shared resource

## Why this helps?

- Since the low-priority task gets temporarily boosted priority, it keeps medium priority tasks from pre-empting the (originally) low priority task
- Once resource released, low-priority task continues at its original priority

# Priority Donation Example



Pintos Lab 1 Exercise 2.2

Details in lab 1 overview session

# Combining Algorithms

**Different types of jobs have different preferences**

- Interactive, CPU-bound, batch, system, etc.
- Hard to use one size to fit all

**Combining scheduling algorithms to optimize for multiple objectives**

- Have multiple queues
- Use a different algorithm for each queue
- Move processes among queues

**Example: Multiple-level feedback queues (MLFQ)**

# Multiple-level Feedback Queues (MLFQ)

**Developed by Fernando J. Corbató in 1962**

- Corbató received the 1990 Turing Award for this work and other work in Multics

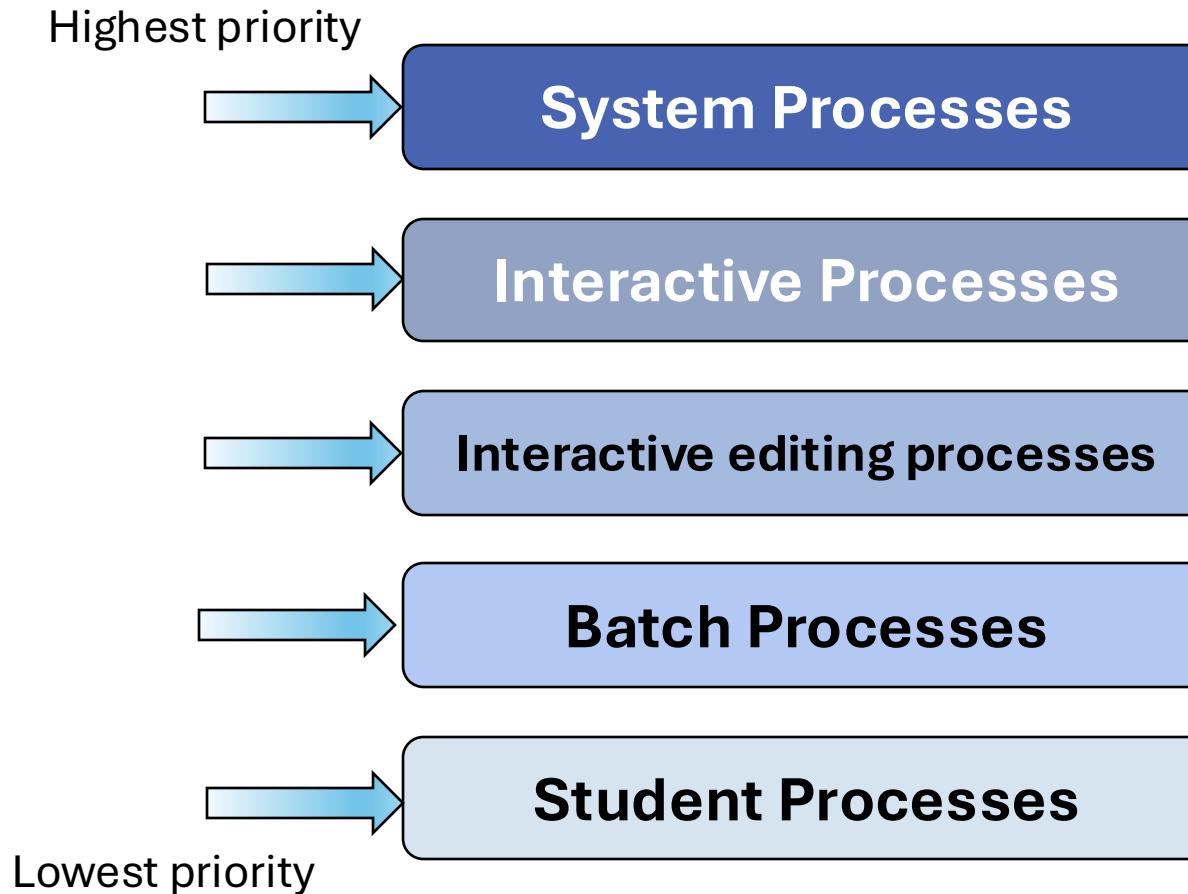
**Widely used in mainstream OSes: Unix, BSD, Windows, MacOS**

**You'll get hands-on experience with it in Lab 1**

**Idea:**

- Multiple queues representing different job types
- Queues w/ priorities: jobs in higher-priority queue preempt jobs lower-priority queue
- Jobs on same queue use the same scheduling algorithm, typically RR

# Multiple-level Queues Scheduling



# Multiple-level Feedback Queues Scheduling

**Goal #1: Optimize job turnaround time for “batch” jobs**

**Goal #2: Minimize response time for “interactive” jobs**

## Challenge:

- No *a priori* knowledge of what type a job is, what the next burst is, etc.
- Let a job tells us its “niceness” (priority)?

## Idea:

- Change a process’s priority based on how it behaves in the past (history “feedback”)

# How to Change Priority Over Time

## Attempt

- *Rule A*: Processes start at top priority
- *Rule B*: If job uses whole slice, demote process
  - i.e., longer time slices at lower priorities
- Example : A long-running “batch” job

## Problems:

- starvation
- gaming the system
  - E.g., performing I/O right before time-slice ends

# How to Change Priority Over Time

## Fixing the problems:

- Periodically **boost** priority for jobs that haven't been scheduled
- Account for job's ***total*** run time at priority level (instead of just this time slice)

# MLFQ in BSD

**Every runnable process on one of 32 run queues**

- Kernel runs process on highest-priority non-empty queue
- Round-robin among processes on same queue

**Process priorities dynamically computed**

- Processes moved between queues to reflect priority changes

**Favor interactive jobs that use less CPU**

# Process Priority Calculation in BSD

**p\_estcpu** – per-process estimated CPU usage

**p\_nice** – user-settable weighting factor, value range [-20, 20]

**Process priority p\_usrpri**

$$p_{\text{usrpri}} \leftarrow 50 + \left( \frac{p_{\text{estcpu}}}{4} \right) + 2 \times p_{\text{nice}}$$

- Calculated every 4 ticks, values are bounded to [50, 127]
- Decrease priority linearly based on recent CPU

**How to calculate p\_estcpu ?**

- Incremented whenever timer interrupt found process running
- Decayed every second while process runnable

$$p_{\text{estcpu}} \leftarrow \left( \frac{2 \times \text{load}}{2 \times \text{load} + 1} \right) \times p_{\text{estcpu}} + p_{\text{nice}}$$

- Load is sampled average of length of run queue plus short-term sleep queue over last minute

# Tips for Pintos

## Same basic idea for second half of Lab 1

- But 64 priorities, not 128
- Higher numbers mean higher priority (in BSD, higher numbers means lower priority)
- Okay to have only one run queue if you prefer (less efficient, but we won't deduct points for it)

## Have to negate priority equation:

In BSD       $p_{usrpri} \leftarrow 50 + \left( \frac{p_{estcpu}}{4} \right) + 2 \times p_{nice}$

In Pintos     $p_{usrpri} \leftarrow 63 + \left( \frac{recent\_cpu}{4} \right) + 2 \times nice$

# **Scheduling Summary**

**Scheduling algorithm determines which process runs, quantum, priority...**

**Many potential goals of scheduling algorithms**

- Utilization, throughput, wait time, response time, etc.

**Various algorithms to meet these goals**

- FCFS/FIFO, SJF, RR, Priority

**Can combine algorithms**

- Multiple-Level Feedback Queues (MLFQ)

# **Next Time**

**Read Chapter 28,29**