

# CE 440 Introduction to Operating System

## Lecture 7: Semaphores and Monitors Fall 2025

Prof. Yigong Hu

BOSTON  
UNIVERSITY

Slides courtesy of Manuel Egele, Ryan Huang and Baris Kasikci

# **Administrivia**

## **Project group member sign up due this Sunday**

- Due this Sunday

## **Lab 1 overview session this Friday**

- 2:30 - 4:00 PM, PHO305

# Recap: Synchronization

**Problem: concurrent threads accessed a **shared resource** without any **synchronization****

- Known as **a race condition**

**The execution of the two threads can be interleaved**

```
balance = get_balance(account);  
balance = balance - amount;
```

```
balance = get_balance(account);  
balance = balance - amount;  
put_balance(account, balance);
```

```
put_balance(account, balance);
```



# Recap: How to Protect Shared Resource?

## 1. Mutual exclusion (mutex)

- If one thread is in the critical section, then no other is

## 2. Progress

- If some thread T is not in the critical section, then T cannot prevent some other thread S from entering the critical section
- A thread in the critical section will eventually leave it

## 3. Bounded waiting (no starvation)

- If some thread T is waiting on the critical section, then T will eventually enter the critical section

# Recap: How to Protect Shared Resource?

## 4. Performance

- The overhead of entering and exiting the critical section is small with respect to the work being done within it

## In summary:

- **Safety property**: nothing bad happens
  - Mutex
- **Liveness property**: something good happens
  - Progress, Bounded Waiting
- **Performance requirement**
  - Performance

**Note: correctness of concurrent is guarantee by design**

# Recap: Lock

**Code that uses mutual exclusion to synchronize its execution is called a critical section**

**A lock is an object in memory providing two operations**

- `acquire()`: wait until lock is free, then take it to enter a C.S
- `release()`: release lock to leave a C.S, waking up anyone waiting for it

# Recap: Higher-Level Synchronization

We looked at using locks to provide **mutual exclusion**

**Locks work, but they have limited semantics**

- Just provide mutual exclusion
- Wasteful

**Instead, we need synchronization mechanisms that**

- Block waiters, leave interrupts enabled in critical sections
- Provide semantics beyond mutual exclusion

**Look at two common high-level mechanisms**

- Semaphores: binary (mutex) and counting
- Monitors: mutexes and condition variables

# Semaphores

**Semaphores have a non-negative integer that supports the two operations:**


- `Semaphore::P()` decrements, blocks until semaphore is open, a.k.a `wait()`
- `Semaphore::V()` increments, allows another thread to enter, a.k.a `signal()`
- **That's it! No other operations – not even just reading its value**
  - Both P and V are after the Dutch word “Proberen” (to try), “Verhogen” (increment)

**Semaphore safety property: the semaphore value is always greater than or equal to 0**

# Using Semaphores to Fix Banking Problem

Use is similar to our locks, but semantics are different

```
struct Semaphore {  
    int value;  
    Queue q;  
} S;  
  
withdraw (account, amount) {  
    P(S);  
    balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    v(S);  
    return balance;  
}
```



```
P(S);  
balance = get_balance(account);  
balance = balance - amount;  
  
P(S);  
  
P(S);  
put_balance(account, balance);  
v(s);  
  
...  
v(s);  
  
...  
v(s);
```

# Semaphores

**Semaphores have a non-negative integer that supports the two operations:**

- `Semaphore::P()` decrements, blocks until semaphore is open, a.k.a `wait()`
- `Semaphore::V()` increments, allows another thread to enter, a.k.a `signal()`
- **That's it! No other operations – not even just reading its value**
  - Both P and V are after the Dutch word “Proberen” (to try), “Verhogen” (increment)

**Semaphore safety property: the semaphore value is always greater than or equal to 0**

**Semaphores are a kind of generalized lock**

- First defined by Dijkstra in the “THE” system in 1968
- Main synchronization primitive used in original UNIX

# Semaphores Implementation

Associated with each semaphore is a **queue of waiting threads**

**When P() is called by a thread:**

- If semaphore is **open**, thread continues
- If semaphore is **closed**, thread blocks on queue

**Then V() opens the semaphore:**

- If a thread is waiting on the queue, the thread is unblocked
- If no threads are waiting on the queue, the signal is remembered for the next thread
  - In other words, V() has “history” (c.f., condition vars later)
  - This “history” is a counter

# Recall: Implementing Locks (4)

Block waiters, interrupts enabled in critical sections

```
struct lock {  
    int held = 0;  
    queue Q;  
}  
void acquire (lock) {  
    Disable interrupts;  
    while (lock→held) {  
        put current thread on lock Q;  
        block current thread;  
    }  
    lock→held = 1;  
    Enable interrupts;  
}
```

```
void release (lock) {  
    Disable interrupts;  
    if (Q) remove waiting thread;  
    unblock waiting thread;  
    lock→held = 0  
    Enable interrupts;  
}
```

Pintos [threads/synch.c](#): sema\_down/up

acquire(lock)  
...  
*Critical section*  
...  
release(lock)

Interrupts Disabled  
Interrupts Enabled  
Interrupts Disabled

# Semaphore Types

Semaphores come in two types

## **Mutex semaphore (or binary semaphore)**

- Represents single access to a resource
- Guarantees mutual exclusion to a critical section


## **Counting semaphore (or general semaphore)**

- Represents a resource with many units available, or a resource that allows certain kinds of unsynchronized concurrent access (e.g., reading)
- Multiple threads can pass the semaphore
- Number of threads determined by the semaphore “count”
  - mutex has count = 1, counting has count = N

# Readers/Writers Problem

Consider a shared database


- Two classes of users:
  - Readers – never modify database
  - Writers – read and modify database
- Is using a single lock on the whole database sufficient?
  - Like to have many readers at the same time
  - Only one writer at a time



# Readers/Writers Problem


## Readers/Writers Problem:

- An object is shared among several threads
- Some threads only read the object, others only write it
- How do we control the access pattern?




# Readers/Writers Problem

If we have **multiple readers**




# Readers/Writers Problem

If we have **multiple readers** and **one writer**



# Readers/Writers Problem

If we have **multiple writers**



# Readers/Writers Problem

## Readers/Writers Problem:

- An object is shared among several threads
- Some threads only read the object, others only write it
- We can allow **multiple readers** but only **one writer**
  - Let  $r$  be the number of readers,  $w$  be the number of writers
  - **Safety:**  $(r \geq 0) \wedge (0 \leq w \leq 1) \wedge ((r > 0) \Rightarrow (w = 0))$

**How can we use semaphores to implement this protocol?**

# Basic Readers/Writers Solution

## Safety Constraints:

- Safety:  $(r \geq 0) \wedge (0 \leq w \leq 1) \wedge ((r > 0) \Rightarrow (w = 0))$

## Basic structure of a solution:

- Reader()
  - Wait until no writers
  - Access database
  - Check out – wake up a waiting writer
- Writer()
  - Wait until no active readers or writers
  - Access database
  - Check out – wake up waiting readers or writer

## Start with...

- Semaphore `w_or_r` – exclusive writing or reading

# Using Semaphores for Readers/Writers

w\_or\_r provides mutex between readers and writers

- writer wait/signal, reader wait/signal when readcount goes from 0 to 1 or from 1 to 0

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Readers/Writers Notes

Consider the following sequence of operators:

- W1, R3, R4

**Why do readers use mutex?**

**Why don't writers use mutex?**

**What if the signal() is above “if (readcount == 1)”?**

# Simulation of Readers/Writers Solution

W1 comes first, R1, R2 come along  
w\_or\_r = 1, mutex = 1, readcount = 0

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 comes first, R1, R2 come along  
w\_or\_r = 0, mutex = 1, readcount = 0

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 comes first, R1, R2 come along  
w\_or\_r = 0, mutex = 0, readcount = 0

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 comes first, R1, R2 come along

w\_or\_r = 0, mutex = 0, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 comes first, R1, R2 come along

w\_or\_r = 0, mutex = 0, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 comes first, R1, R2 come along

w\_or\_r = 0, mutex = 0, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue

w\_or\_r = 1, mutex = 0, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue

w\_or\_r = 0, mutex = 0, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue

w\_or\_r = 0, mutex = 1, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue

w\_or\_r = 0, mutex = 0, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue

w\_or\_r = 0, mutex = 0, readcount = 2

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue

w\_or\_r = 0, mutex = 0, readcount = 2

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue


w\_or\_r = 0, mutex = 1, readcount = 2

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read; 
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue

w\_or\_r = 0, mutex = 0, readcount = 2

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue

w\_or\_r = 0, mutex = 0, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue

w\_or\_r = 0, mutex = 0, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue

w\_or\_r = 0, mutex = 0, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue

w\_or\_r = 0, mutex = 0, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue

w\_or\_r = 0, mutex = 1, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue

w\_or\_r = 0, mutex = 0, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue

w\_or\_r = 0, mutex = 0, readcount = 0

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue

w\_or\_r = 0, mutex = 0, readcount = 0

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue

w\_or\_r = 1, mutex = 0, readcount = 0

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue

w\_or\_r = 1, mutex = 0, readcount = 0

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Readers/Writers Notes

Consider the following sequence of operators:

- W1, R3, R4

Why do readers use mutex?

Why don't writers use mutex?

What if the signal() is above “if (readcount == 1)”?

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    signal(&mutex); // unlock readcount
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    Read;
```

# Simulation of Readers/Writers Solution

W1 finishes, R1, R2 continue

w\_or\_r = 1, mutex = 1, readcount = 0

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

Is this safe?

# Readers/Writers Notes

## Is It Safe?

- Yes

If readers and writers are waiting, who goes first?

# If Writer and Reader Are Waiting for Writer

W1 comes first, R1, W2 come along

w\_or\_r = 1, mutex = 1, readcount = 0

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# If Writer and Reader Are Waiting for Writer

W1 comes first, R1, W2 come along

w\_or\_r = 0, mutex = 1, readcount = 0

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# If Writer and Reader Are Waiting for Writer

W1 comes first, R1, W2 come along

w\_or\_r = 0, mutex = 1, readcount = 0

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# If Writer and Reader Are Waiting for Writer

W1 comes first, R1, W2 come along

w\_or\_r = 0, mutex = 0, readcount = 0

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# If Writer and Reader Are Waiting for Writer

W1 comes first, R1, W2 come along

w\_or\_r = 0, mutex = 0, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# If Writer and Reader Are Waiting for Writer

W1 comes first, R1, W2 come along

w\_or\_r = 0, mutex = 0, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# If Writer and Reader Are Waiting for Writer

W1 comes first, R1, W2 come along

w\_or\_r = 0, mutex = 0, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

Who go first?

# Readers/Writers Notes

If a writer is writing, where will readers be waiting?

- Yes

If readers and writers are waiting, who goes first?

- If waiting for writers, once a writer exits, all readers/writers can fall through
  - Which reader gets to go first?

# If Writer and Reader Are Waiting for Reader

R1 comes first, W1, R2 come along

w\_or\_r = 1, mutex = 1, readcount = 0

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# If Writer and Reader Are Waiting for Reader

R1 comes first, W1, R2 come along

w\_or\_r = 1, mutex = 0, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# If Writer and Reader Are Waiting for Reader

R1 comes first, W1, R2 come along

w\_or\_r = 1, mutex = 0, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# If Writer and Reader Are Waiting for Reader

R1 comes first, W1, R2 come along

w\_or\_r = 0, mutex =0, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# If Writer and Reader Are Waiting for Reader

R1 comes first, W1, R2 come along

w\_or\_r = 0, mutex =0, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

Who go first?

# Readers Always Go First

R1 comes first, W1, R2 come along  
w\_or\_r = 0, mutex =1, readcount = 1

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Readers Always Go First

R1 comes first, W1, R2 come along

w\_or\_r = 0, mutex =0, readcount = 2

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Readers Always Go First

R1 comes first, W1, R2 come along

w\_or\_r = 0, mutex =1, readcount = 2

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Readers Always Go First

R1 comes first, W1, R2 come along

w\_or\_r = 0, mutex =1, readcount = 2

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

# Readers Always Go First

R1 finishes, R2 blocks W1

w\_or\_r = 0, mutex =1, readcount = 2

```
// exclusive writer or reader
Semaphore w_or_r(1);

// number of readers
int readcount = 0;

// mutual exclusion to readcount
Semaphore mutex(1);

writer() {
    wait(&w_or_r); // lock out others
    Write;
    signal(&w_or_r); // up for grabs
}
```

```
reader() {
    wait(&mutex); // lock readcount
    readcount += 1; // one more reader
    if (readcount == 1)
        wait(&w_or_r); // synch w/ writers
    signal(&mutex); // unlock readcount
    Read;
    wait(&mutex); // lock readcount
    readcount -= 1; // one less reader
    if (readcount == 0)
        signal(&w_or_r); // up for grabs
    signal(&mutex); // unlock readcount
}
```

What if we have R3,R4,R5... coming now?

# Readers/Writers Notes

If a writer is writing, where will readers be waiting?

- Yes

If readers and writers are waiting, who goes first?

- If waiting for writers, once a writer exits, all readers/writers can fall through
  - Which reader gets to go first?

# Readers/Writers Notes

If a writer is writing, where will readers be waiting?

- Yes

If readers and writers are waiting, who goes first?

- If waiting for writers, once a writer exits, all readers/writers can fall through
  - Which reader gets to go first?
- If waiting for readers, possible **starvation** for the writer

# Semaphore Questions

**Are there any problems that can be solved with counting semaphores that cannot be solved with mutex semaphores?**

- If a system only gives you mutex semaphore, can you use it to implement counting semaphores?

**Does it matter which thread is unblocked by a signal operation?**

# Tips for Pintos: Semaphore Implementation

```
void sema_down(struct semaphore *sema)
{
    enum intr_level old_level;
    old_level = intr_disable();
    while (sema->value == 0) {
        list_push_back(&sema->waiters,
                       &thread_current()->elem);
        thread_block();
    }
    sema->value--;
    intr_set_level(old_level);
}
```

```
void sema_up(struct semaphore *sema)
{
    enum intr_level old_level;
    old_level = intr_disable();
    if (!list_empty(&sema->waiters))
        thread_unblock(list_entry(
            list_pop_front(&sema->waiters),
            struct thread, elem));
    sema->value++;
    intr_set_level(old_level);
}
```

To reference current thread: `thread_current()`

`thread_block()` puts the current thread to sleep

Lab 1 note:

- leverage semaphore instead of directly using `thread_block()`

# Tips for Pintos: `thread_block()`

Pick another  
thread to run

```
/* Puts the current thread to sleep. This
function
must be called with interrupts turned off.*/
void thread_block ()
{
    ASSERT (!intr_context ());
    ASSERT (intr_get_level () == INTR_OFF);
    thread_current ()->status = THREAD_BLOCKED;
    schedule ();
}
```

**thread\_block() assumes the interrupts are disabled**

**This means we will have the thread sleep with interrupts disabled**

**Isn't this bad?**

- Shouldn't we only disable interrupts when entering/leaving critical sections but keep interrupts enabled during critical section?

# Interrupts Re-enabled Right After Context Switch

```
thread_yield() {  
    Disable interrupts;  
    add current thread to ready_list;  
    schedule(); // context switch  
    Enable interrupts;  
}
```

```
sema_down() {  
    Disable interrupts;  
    while(value == 0) {  
        add current thread to waiters;  
        thread_block();  
    }  
    value--;  
    Enable interrupts;  
}
```

[**thread\_yield**]  
*Disable interrupts;*  
add current thread to ready\_list;  
schedule();

Thread 1

[**thread\_yield**]  
*(Returns from schedule())*  
**Enable interrupts;**

Thread 2

...

[**sema\_down**]  
*Disable interrupts;*  
while(value == 0) {  
 add current thread to waiters;  
 thread\_block();  
}

Thread 2

[**thread\_yield**]  
*(Returns from schedule())*  
**Enable interrupts;**

Thread 1

# Semaphore Summary

**Semaphores can be used to solve any traditional sync. Problems**

**However, they have some drawbacks**

- They are essentially shared global variables
  - Can potentially be accessed anywhere in program
- No connection between the semaphore and the data controlled by the semaphore
- Used both for critical sections (mutual exclusion) and coordination (scheduling)
  - Note that I had to use comments in the code to distinguish
- No control or guarantee of proper usage

**Sometimes hard to use and prone to bugs**

# Semaphores are good but... Monitors are better!

Semaphores are a huge step up; just think of trying to do the reader/writer with only loads and stores or lock

**Problem is that semaphores are dual purpose:**

- They are used for both mutex and scheduling constraints

**Insight:**

- Use **locks** for mutual exclusion and **condition variables** for scheduling constraints
- Use programming language support

# Monitor

**A programming language construct that controls access to shared data**

- Synchronization code added by compiler, enforced at runtime
- Why is this an advantage?

**A monitor is a module that encapsulates**

- Shared data structures
- Procedures that operate on the shared data structures
- Synchronization between concurrent threads that invoke the procedures

```
Monitor account {  
    double balance;  
  
    double withdraw (amount) {  
        balance = balance - amount;  
        put_balance(account, balance);  
        return balance;  
    }  
}
```

# Monitor

**A programming language construct that controls access to shared data**

- Synchronization code added by compiler, enforced at runtime
- Why is this an advantage?

**A monitor is a module that encapsulates**

- Shared data structures
- Procedures that operate on the shared data structures
- Synchronization between concurrent threads that invoke the procedures

**A monitor protects its data from unstructured access**

**It guarantees that threads accessing its data through its procedures interact only in legitimate ways**

# Bank Account Problem With Monitor

```
Monitor account {  
    double balance;  
  
    double withdraw (amount) {  
        balance = balance - amount;  
        put_balance(account, balance);  
        return balance;  
    }  
}
```

Threads  
block

When first thread exits, another can  
enter. Which one is undefined

```
withdraw(amount)  
balance = balance - amount;
```

```
withdraw(amount)
```

```
withdraw(amount)
```

```
return balance
```

```
balance = balance - amount;  
return balance;
```

```
balance = balance - amount;  
return balance;
```

# Monitor Semantics

A monitor guarantees **mutual exclusion**

- Only one thread can execute any monitor procedure at any time
  - The thread is “in the monitor”
- If a second thread invokes a monitor procedure when a first thread is already executing one, it blocks
  - So the monitor has to have a wait queue...
- If a thread within a monitor blocks, another one can enter

What are the implications in terms of parallelism in a monitor?

A **monitor invariant** is a **safety property** associated with the monitor

- It's expressed over the monitored variables.
- It holds whenever a thread enters or exits the monitor.

# Condition Variables

**But what if a thread wants to wait for something inside the monitor?**

- If we busy wait, it's bad
- Even worse, no one can get in the monitor to make changes now!

A **condition variable** is associated with a **condition** needed for a thread to make progress once it is in the monitor.

```
Monitor M {  
    ... monitored variables  
    Condition c;  
  
    void enterMonitor (...) {  
        if (extra property not true) wait(c); waits outside of the monitor's mutex  
        do what you have to do  
        if (extra property true) signal(c); brings in one thread waiting on condition  
    }
```

# Condition Variables

**Condition variables support three operations:**

- **Wait** – release monitor lock, wait for C/V to be signaled
  - So condition variables have wait queues, too
- **Signal** – wakeup one waiting thread
- **Broadcast** – wakeup all waiting threads

**Condition variables are not boolean objects**

- `if (condition_variable) then...` does not make sense
- `if (num_resources == 0) then wait(resources_available)` does
- We will explain the detail in next lecture

# Condition Variables != Semaphores

## Condition variables != semaphores

- Although their operations have the same names, they have entirely different semantics (such is life, worse yet to come)
- However, they each can be used to implement the other

## Access to the monitor is controlled by a lock

- `wait()` blocks the calling thread, and **gives up the lock**
  - To call `wait`, the thread has to be in the monitor (hence has lock)
  - `Semaphore::wait` just blocks the thread on the queue
- `signal()` causes a waiting thread to wake up
  - **If there is no waiting thread, the signal is lost**
  - `Semaphore::signal` increases the semaphore count, allowing future entry even if no thread is waiting
  - Condition variables have no history

# Signal Semantics

Two flavors of monitors that differ in the scheduling semantics of signal()

- Hoare monitors (original)
  - signal() immediately switches from the caller to a waiting thread
  - The condition that the waiter was anticipating is guaranteed to hold when waiter executes
  - Signaler must restore monitor invariants before signaling

## Hoare

```
if (!condition)
    wait(cond_var);
```

Condition definitely holds since we  
just context switched from signal

# Signal Semantics

- Mesa monitors (Mesa, Java)
  - signal() places a waiter on the ready queue, **but signaler continues inside monitor**
  - Condition is not necessarily true when waiter runs again
    - Returning from wait() is only a *hint* that something changed
    - Must recheck conditional case

## Mesa

```
while (!condition)
    wait(cond_var); condition might have been changed, if so, wait again
condition holds now
```

# Hoare vs. Mesa Monitors

## Tradeoffs

- Mesa monitors easier to use, more efficient
  - Fewer context switches, easy to support broadcast
- Hoare monitors leave less to chance
  - Easier to reason about the program

# Summary

## Semaphores

- `wait()`/`signal()` implement blocking mutual exclusion
- Also used as atomic counters (counting semaphores)
- Can be inconvenient to use

## Monitors

- Synchronizes execution within procedures that manipulate encapsulated data shared among procedures
  - Only one thread can execute within a monitor at a time
- Relies upon high-level language support

## Condition variables

- Used by threads as a synchronization point to wait for events
- Inside monitors, or outside with locks