

# CE 440 Introduction to Operating System

## Lecture 12: Page Replacement Fall 2025

Prof. Yigong Hu



Slides courtesy of Manuel Egele, Ryan Huang and Baris Kasikci

# **Administrivia**

- Next next Wednesday (October 22nd) is Midterm
- Midterm Review Session on next Monday

# Memory Management

Next four lectures are going to cover memory management

Goals of memory management

Mechanisms

- Physical and virtual addressing (1)
- Techniques: partitioning, paging, segmentation (1)
- Page table management, TLBs, VM tricks (2)

Policies

- Page replacement algorithms (3)

# Lecture Overview

**Review paging and page replacement**

**Page replacement algorithms**

**Local vs. global replacement**

**Thrashing**

# Recap: Demand Paging

## Paging (swapping) from the OS perspective:

- Pages are evicted to disk when memory is full
- Pages loaded from disk when referenced again
- References to evicted pages cause a TLB miss
  - PTE was invalid, causes fault
- OS allocates a page frame, reads page from disk to the page frame
- OS fills in PTE, marks it valid, and restarts faulting instruction

## Dirty vs. clean pages

- Actually, only dirty pages (modified) need to be written to disk
- Clean pages do not
  - but you do need to know where on disk to read them from again!

# Paging Challenges

## How to resume a process after a fault?

- Need to save state and resume

## What to fetch from disk?

- Just needed page or more?

## What to evict?

- How to allocate physical pages amongst processes?
- Which of a particular process's pages to keep in memory?
- A poor choice can lead to horrible performance
  - cost of paging: disk much, much slower than memory

# Page Replacement

**When a page fault occurs, the OS loads the faulted page from disk into a page frame of physical memory**

**At some point, the process used all of the page frames it is allowed to use**

- This is likely (much) less than all of available memory

**When this happens, the OS must **replace** a page for each page faulted in**

- It must evict a page to free up a page frame

**The **page replacement algorithm** determines how this is done**

- Greatly affect performance of paging (virtual memory)
- Also called page eviction policies

# Locality

**All paging schemes depend on locality**

- Processes reference pages in localized patterns

**Temporal locality**

- Locations referenced recently likely to be referenced again

**Spatial locality**

- Locations near recently referenced locations are likely to be referenced soon

**While the cost of paging is high, if it is infrequent enough it is acceptable**

- Processes usually exhibit both kinds of locality during a run, making paging practical

# Lecture Overview

Review paging and page replacement

**Page replacement algorithms**

**Local vs. global replacement**

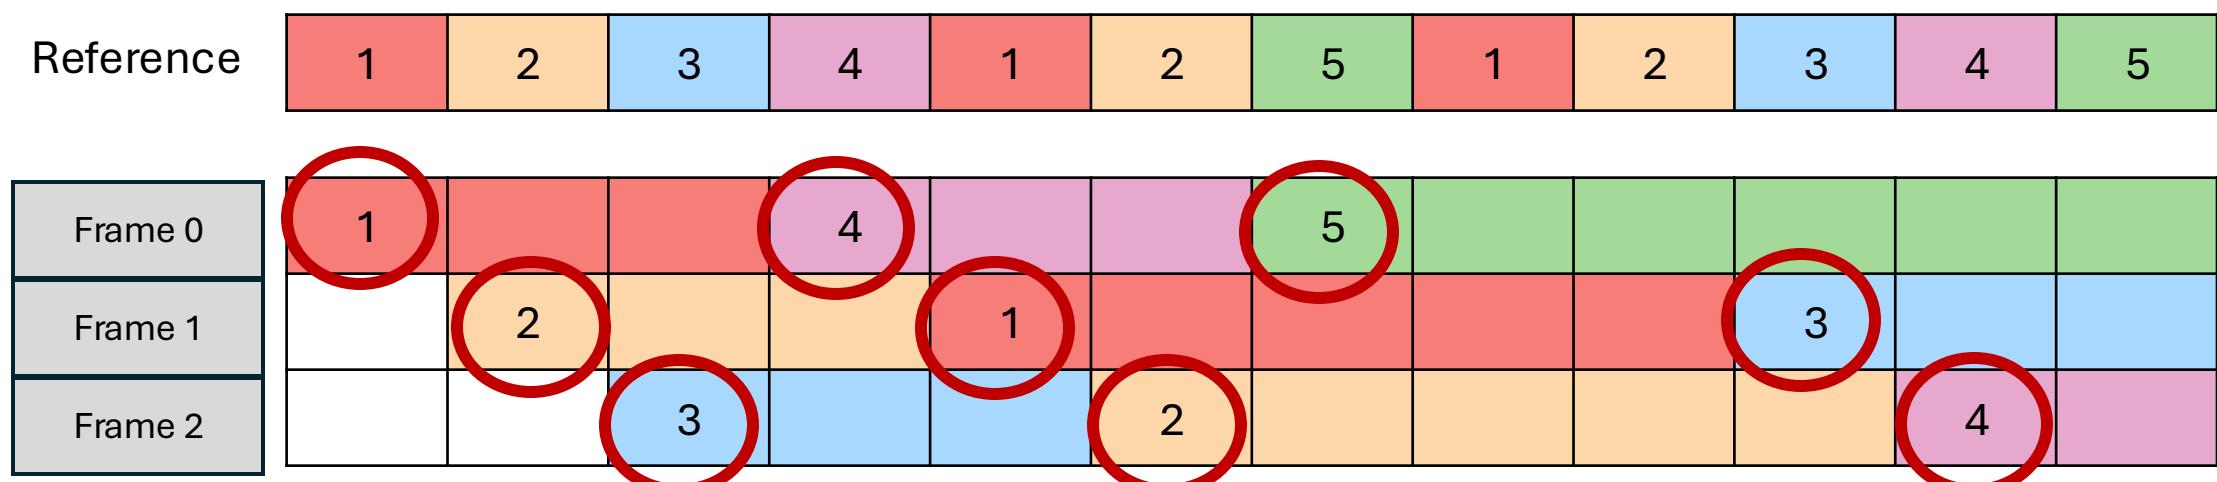
**Thrashing**

# First-In First-Out (FIFO)

Evict oldest fetched page in system

Example: suppose we have 3 page frames, 5 virtual pages and the following reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

FIFO Page replacement: 9 page faults



# First-In First-Out (FIFO)

Evict oldest fetched page in system

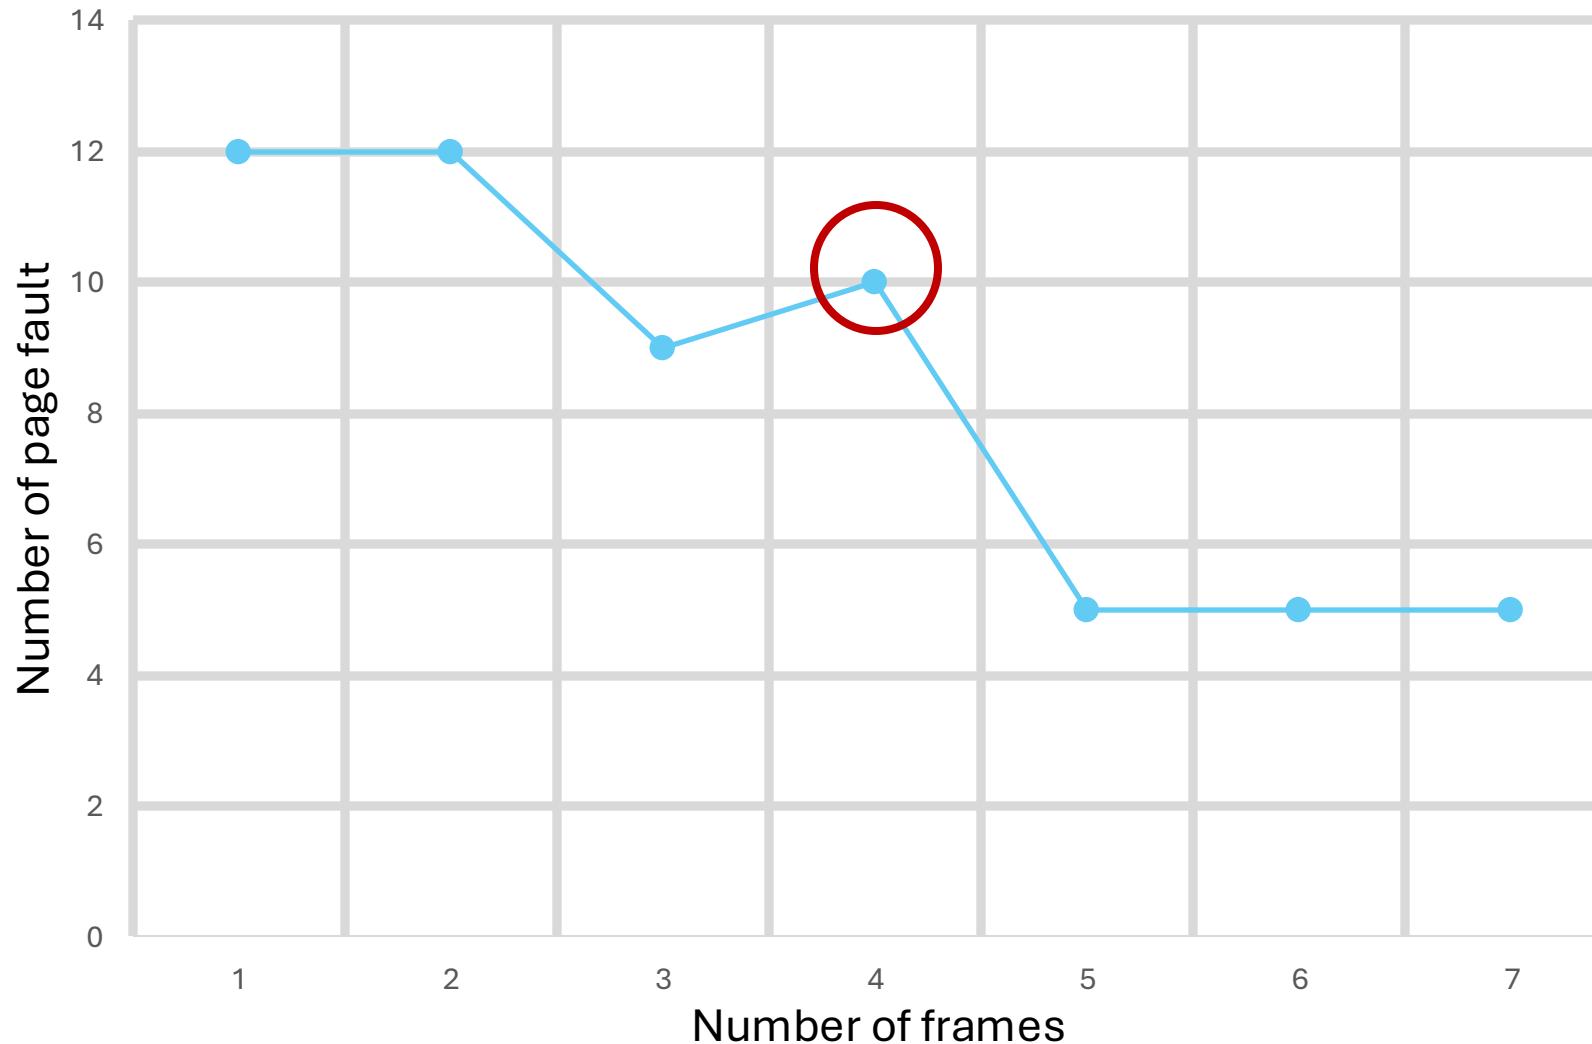
Example: suppose we have **4 page frames**, 5 virtual pages and the following reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

FIFO Page replacement: 10 page faults

Reference	1	2	3	4	1	2	5	1	2	3	4	5
Frame 0	1						5			4		
Frame 1		2						1			5	
Frame 2			3						2			
Frame 4				4						3		

The diagram illustrates the FIFO page replacement algorithm. It shows a reference string of 12 pages (1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5) at the top. Below it, four frames are managed over 12 time steps. The frames are initially empty (white). The sequence of page arrivals is: Frame 0 gets page 1; Frame 1 gets page 2; Frame 2 gets page 3; Frame 4 gets page 4. In the next step, page 1 arrives again and is placed in Frame 0, causing page 2 to be evicted (circled in red). This pattern repeats: page 2 arrives in Frame 1, causing page 3 to be evicted (circled in red); page 5 arrives in Frame 0, causing page 1 to be evicted (circled in red); page 1 arrives in Frame 1, causing page 2 to be evicted (circled in red); page 2 arrives in Frame 2, causing page 3 to be evicted (circled in red); page 3 arrives in Frame 4, causing page 4 to be evicted (circled in red); page 4 arrives in Frame 0, causing page 5 to be evicted (circled in red); page 5 arrives in Frame 1, causing page 1 to be evicted (circled in red); page 1 arrives in Frame 2, causing page 2 to be evicted (circled in red); page 2 arrives in Frame 4, causing page 3 to be evicted (circled in red); page 3 arrives in Frame 0, causing page 4 to be evicted (circled in red); page 4 arrives in Frame 1, causing page 5 to be evicted (circled in red). The final state shows Frame 0 with page 5, Frame 1 with page 4, Frame 2 with page 5, and Frame 4 with page 3.

# Belady's Anomaly



**More physical memory doesn't always mean fewer faults**

# Optimal Page Replacement

**What is optimal (if you knew the future)?**

- Replace page that will not be used for longest period of time

Example: suppose we have **4 page frames**, 5 virtual pages and the same reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

Optimal Page replacement: 6 page faults

Reference	1	2	3	4	1	2	5	1	2	3	4	5
Frame 0	1										4	
Frame 1		2										
Frame 2			3									
Frame 4				4			5					

The diagram illustrates the Optimal Page Replacement algorithm. The top row shows the reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5. The bottom part shows a memory structure with four frames. The first frame contains page 1, circled in red. The second frame contains page 2, circled in red. The third frame contains page 3, circled in red. The fourth frame contains page 4, circled in red. When page 5 arrives, it is placed in the fourth frame because it has the longest future usage. This results in 6 page faults.

# Belady's Algorithm

**Known as the optimal page replacement algorithm**

- Rationale: the best page to evict is the one never touched again
- Never is a long time, so picking the page closest to “never” is the next best thing
- Proved by Belady

**Problem: Have to predict the future**

**Why is Belady's useful then? Use it as a yardstick**

- Compare page replacement algorithms with the optimal to gauge room for improvement
- If optimal is not much better, then algorithm is pretty good
- If optimal is much better, then algorithm could use some work
  - Random replacement is often the lower bound

# Least Recently Used (LRU)

**Approximate optimal with least recently used**

- Because past often predicts the future
- On replacement, evict the page that has not been used for the longest time in the **past** (Belady's: **future**)

LRU Page replacement: 8 page faults

Reference	1	2	3	4	1	2	5	1	2	3	4	5
Frame 0	1											5
Frame 1		2										
Frame 2			3					5			4	
Frame 4				4					3			

# Least Recently Used (LRU)

## Approximate optimal with least recently used

- Because past often predicts the future
- On replacement, evict the page that has not been used for the longest time in the **past** (Belady's: **future**)

LRU Page replacement: 8 page faults

## Problem 1: **Can be pessimistic – example?**

- Looping over memory (then want MRU eviction)

## Problem 2: How to implement?

# Strawman LRP Implementation

## Stamp PTEs with timer value

- E.g., CPU has cycle counter
- Automatically writes value to PTE on each page access
- Scan page table to find oldest counter value = LRU page
- Problem: Would double memory traffic!

## Keep doubly-linked list of pages

- On access remove page, place at tail of list
- Problem: again, very expensive

## What to do?

- Just approximate LRU, don't try to do it exactly

# Clock Algorithm

**Use accessed bit supported by most hardware**

- E.g., Pentium will write 1 to A bit in PTE on first access
- Software managed TLBs like MIPS can do the same

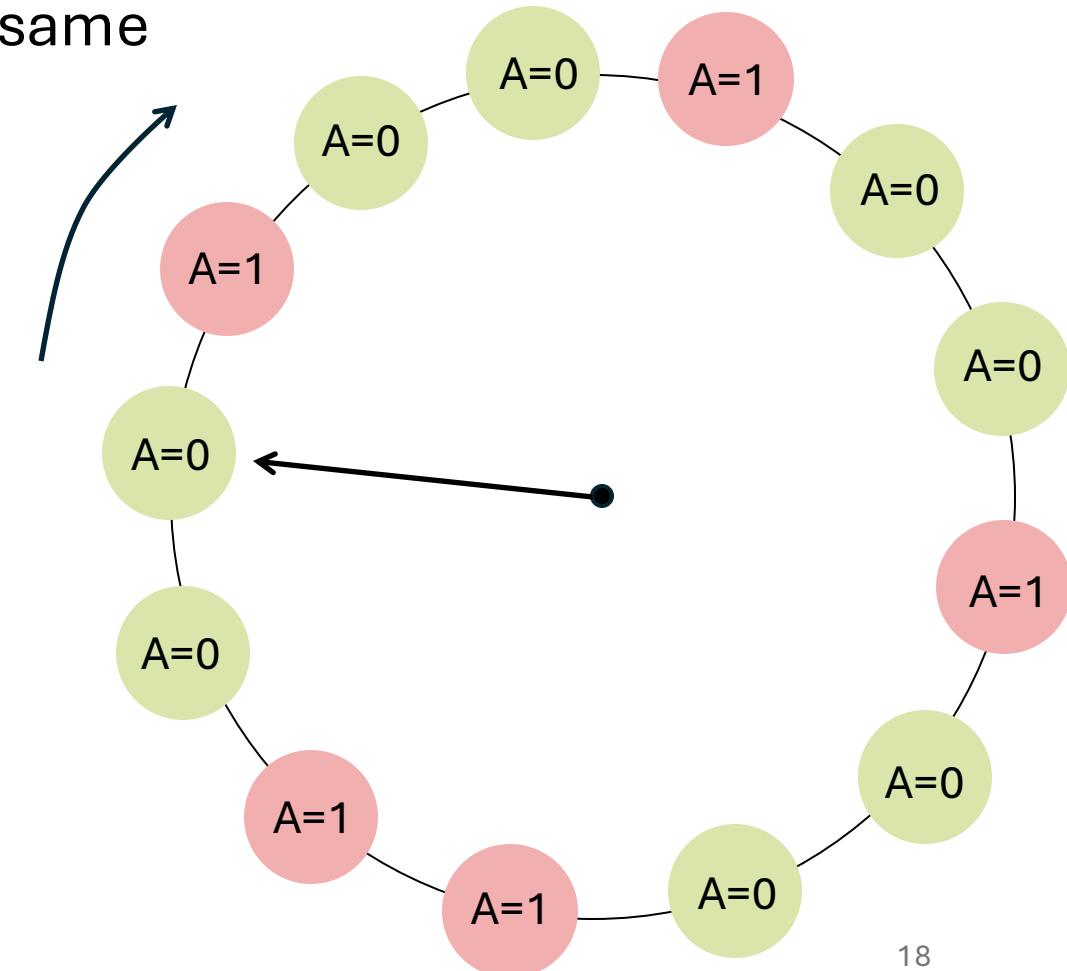
**Do FIFO but skip accessed pages**

**Keep pages in circular FIFO list**

**Scan:**

- page's A bit = 1, set to 0 & skip
- else if A = 0, evict

**A.k.a. second-chance replacement**



# Clock Algorithm

**Use accessed bit supported by most hardware**

- E.g., Pentium will write 1 to A bit in PTE on first access
- Software managed TLBs like MIPS can do the same

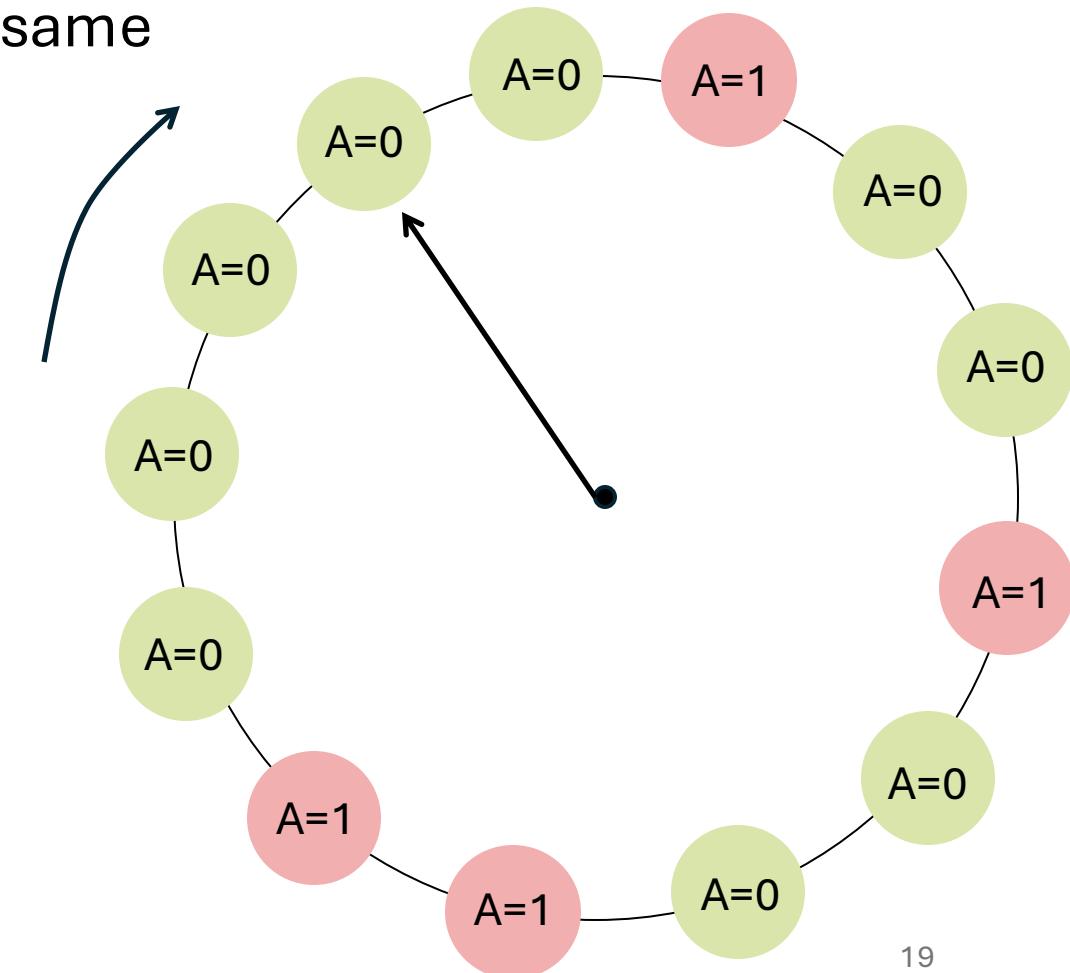
**Do FIFO but skip accessed pages**

**Keep pages in circular FIFO list**

**Scan:**

- page's A bit = 1, set to 0 & skip
- else if A = 0, evict

**A.k.a. second-chance replacement**



# Clock Algorithm

**Use accessed bit supported by most hardware**

- E.g., Pentium will write 1 to A bit in PTE on first access
- Software managed TLBs like MIPS can do the same

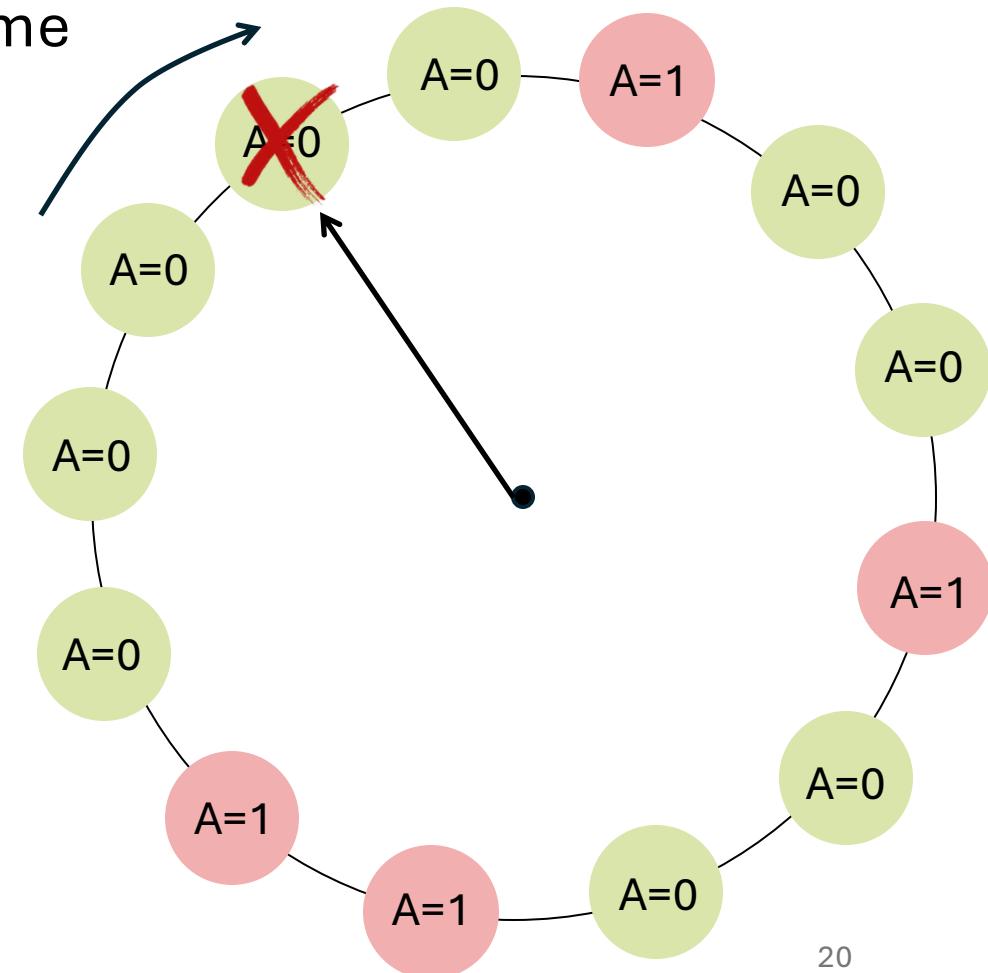
**Do FIFO but skip accessed pages**

**Keep pages in circular FIFO list**

**Scan:**

- page's A bit = 1, set to 0 & skip
- else if A = 0, evict

**A.k.a. second-chance replacement**



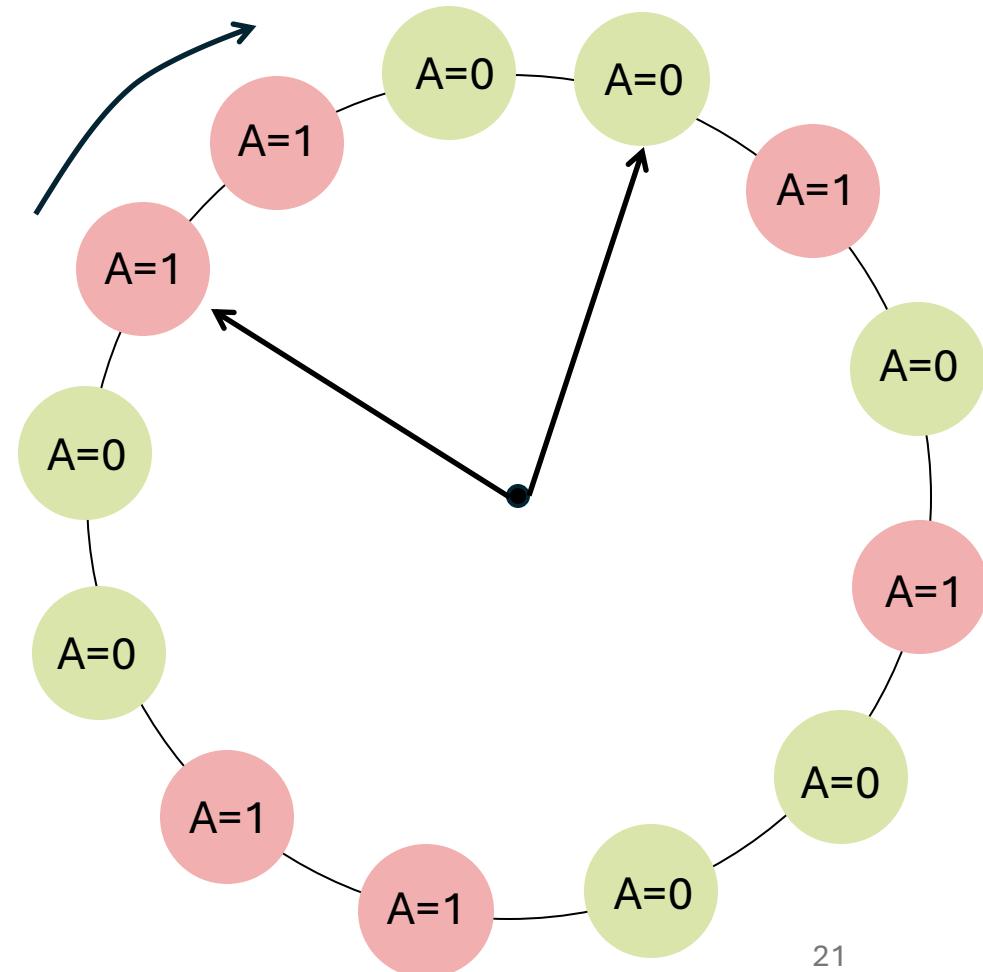
# Clock Algorithm (2)

**Large memory may be a problem**

- Most pages referenced in long interval

**Add a second clock hand**

- Two hands move in lockstep
- Leading hand clears A bits
- Trailing hand evicts pages with A=0



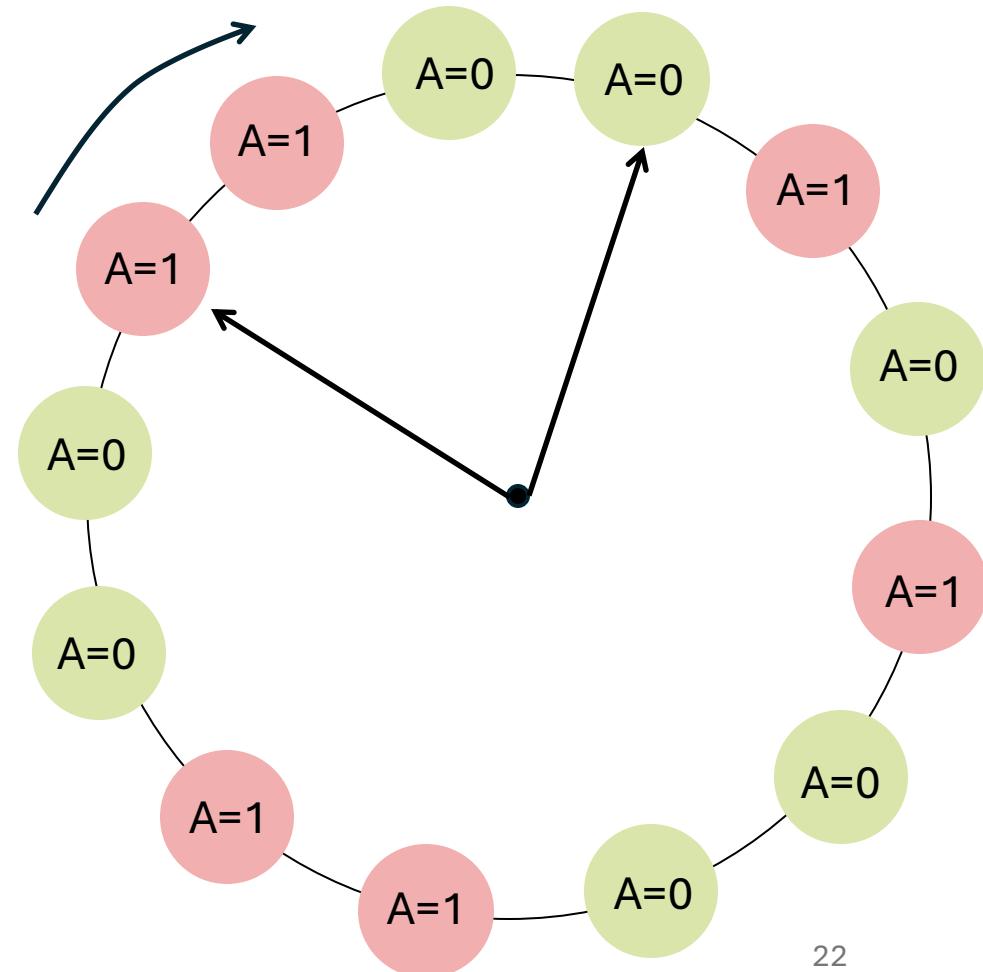
# Clock Algorithm (2)

**Large memory may be a problem**

- Most pages referenced in long interval

**Add a second clock hand**

- Two hands move in lockstep
- Leading hand clears A bits
- Trailing hand evicts pages with A=0



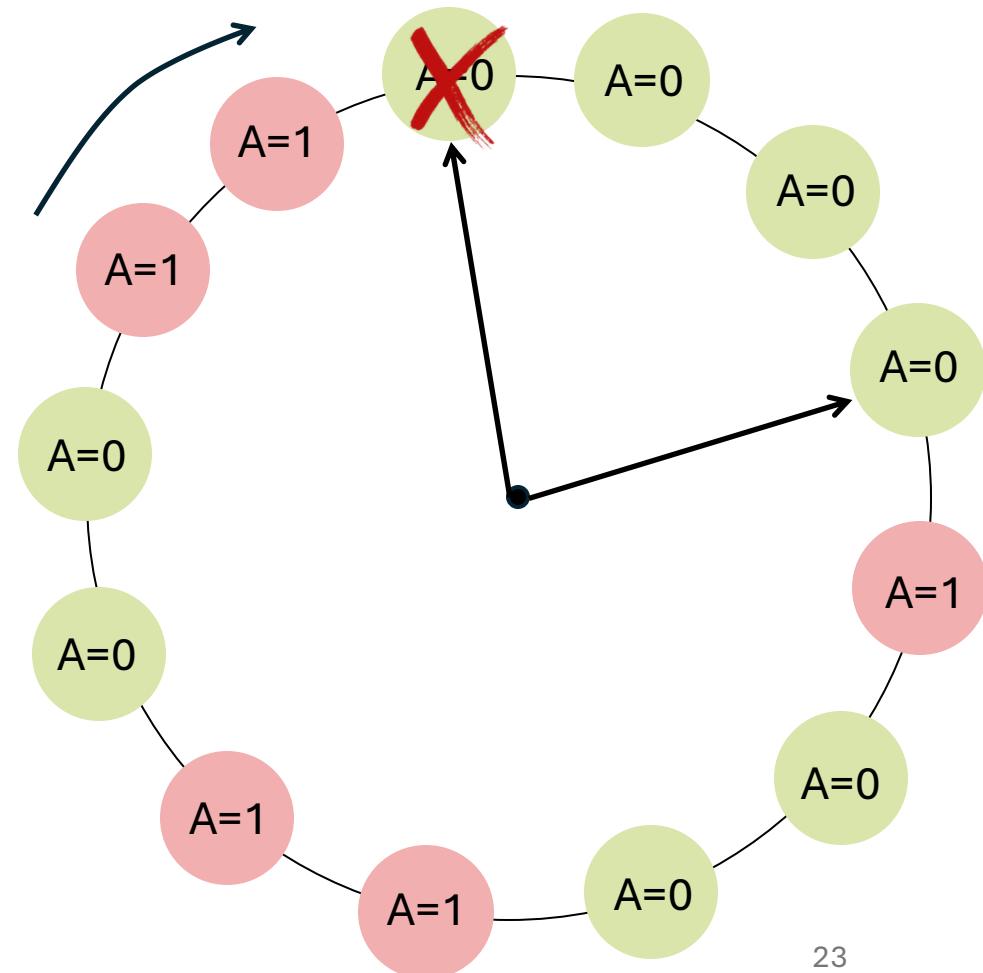
# Clock Algorithm (2)

**Large memory may be a problem**

- Most pages referenced in long interval

**Add a second clock hand**

- Two hands move in lockstep
- Leading hand clears A bits
- Trailing hand evicts pages with A=0



# Other Replacement Algorithms

## Random eviction

- Dirt simple to implement
- Not overly horrible (avoids Belady & pathological cases)

## LFU (least frequently used) eviction

- Instead of just A bit, count # times each page accessed
- Least frequently accessed must not be very useful (or maybe was just brought in and is about to be used)
- Decay usage counts over time (for pages that fall out of usage)

## MFU (most frequently used) algorithm

- Because page with the smallest count was probably just brought in and has yet to be used

## Neither LFU nor MFU used very commonly

# Lecture Overview

Review paging and page replacement

Page replacement algorithms

**Local vs. global replacement**

**Thrashing**

# Fixed vs. Variable Space

**How to determine how much memory to give to each process?**

## Fixed space algorithms

- Each process is given a limit of pages it can use
- When it reaches the limit, it replaces from its own pages
  - Local replacement
- Some processes may do well while others suffer

## Variable space algorithms

- Process' set of pages grows and shrinks dynamically
- Global replacement
  - One process can ruin it for the rest

# Working Set Model

**A working set of a process is used to model the dynamic locality of its memory usage**

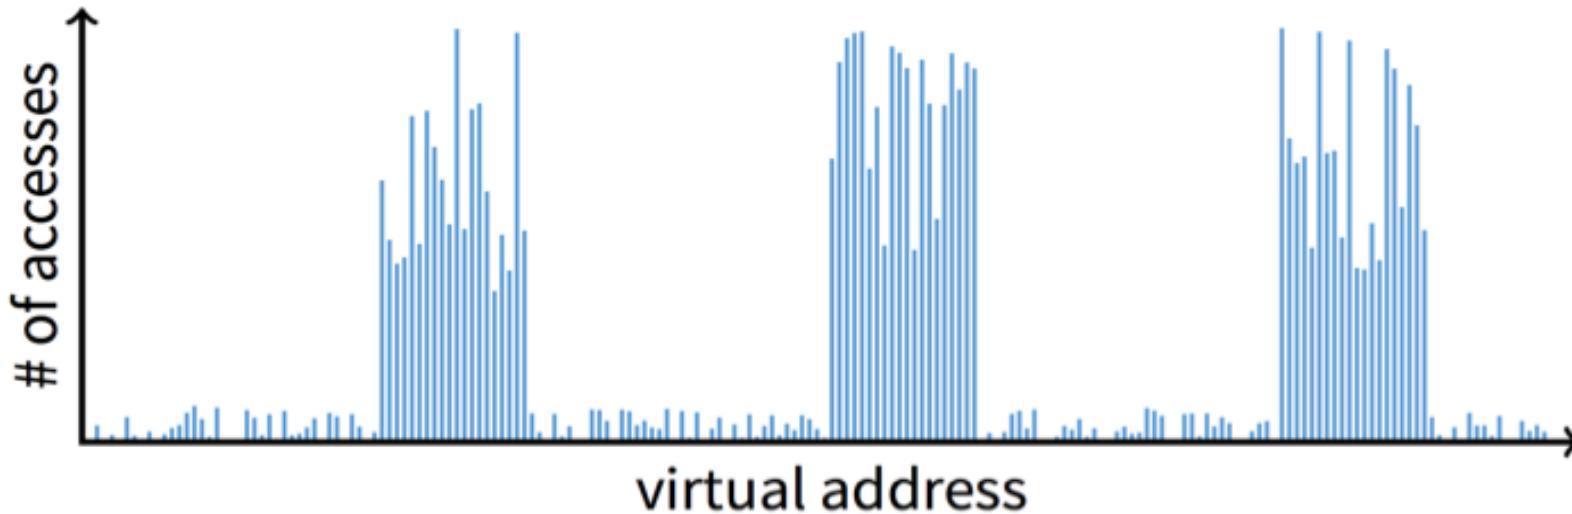
- Defined by Peter Denning in 60s, published at the first SOSP conference

## Definition

- $WS(t, w) = \{ \text{pages } P \text{ such that } P \text{ was referenced in the time interval } (t, t-w) \}$
- $t$  – time,  $w$  – working set window (measured in page refs)

**A page is in the working set (WS) only if it was referenced in the last  $w$  references**

# Working Set Model



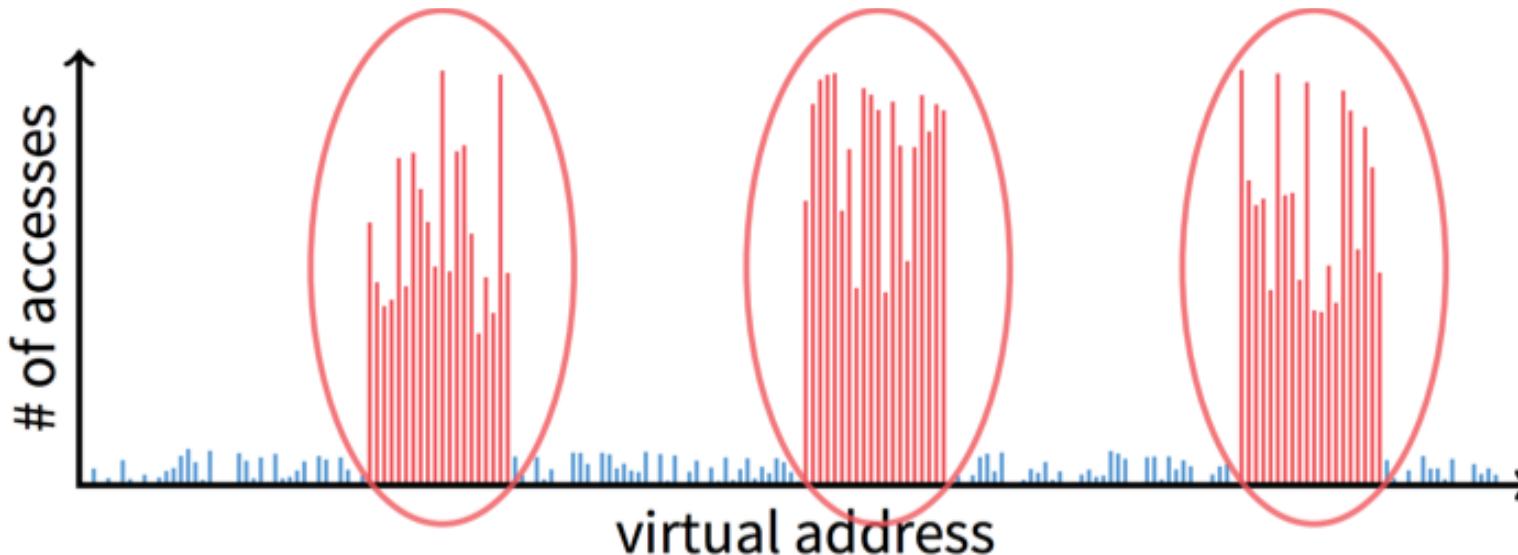
**Disk much, much slower than memory**

- Goal: run at memory speed, not disk speed

**80/20 rule: 20% of memory gets 80% of memory accesses**

- Keep the hot 20% in memory
- Keep the cold 80% on disk

# Working Set Model



**Disk much, much slower than memory**

- Goal: run at memory speed, not disk speed

**80/20 rule: 20% of memory gets 80% of memory accesses**

- Keep the hot 20% in memory
- Keep the cold 80% on disk

# Working Set Model



**Disk much, much slower than memory**

- Goal: run at memory speed, not disk speed

**80/20 rule: 20% of memory gets 80% of memory accesses**

- Keep the hot 20% in memory
- **Keep the cold 80% on disk**

# Working Set Size

**The working set size is the # of unique pages in the working set**

- The number of pages referenced in the interval  $(t, t - w)$

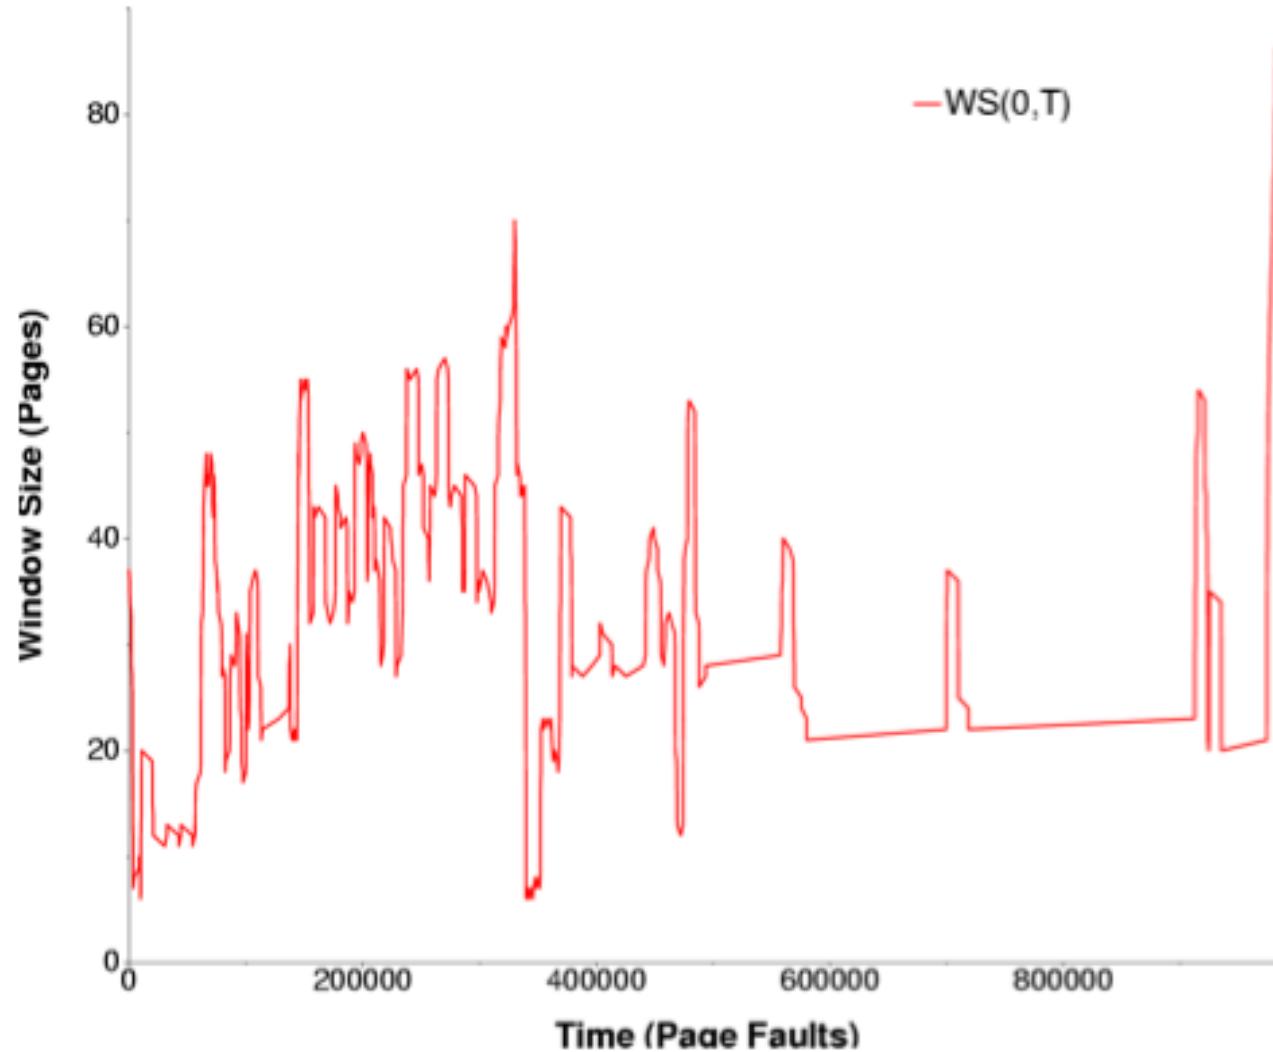
**The working set size changes with program locality**

- During periods of poor locality, you reference more pages
- Within that period of time, the working set size is larger

**Intuitively, want the working set to be the set of pages a process needs in memory to prevent heavy faulting**

- Each process has a param  $w$  that determines a working set with few faults
- Denning: Don't run a process unless working set is in memory

# Example: gcc Working Set



# Working Set Problems

## Problems

- How do we determine w?
- How do we know when the working set changes?

## Too hard to answer

- So, working set is not used in practice as a page replacement algorithm

## However, it is still used as an abstraction

- The intuition is still valid
- When people ask, “How much memory does Firefox need?”, they are in effect asking for the size of Firefox’s working set

# Page Fault Frequency (PFF)

**Page Fault Frequency (PFF) is a variable space algorithm that uses a more ad-hoc approach**

- Monitor the fault rate for each process
- If the fault rate is above a high threshold, give it more memory
  - So that it faults less
  - But not always (FIFO, Belady's Anomaly)
- If the fault rate is below a low threshold, take away memory
  - Should fault more
  - But not always

**Hard to use PFF to distinguish between changes in locality and changes in size of working set**

# Lecture Overview

Review paging and page replacement

Page replacement algorithms

Local vs. global replacement

**Thrashing**

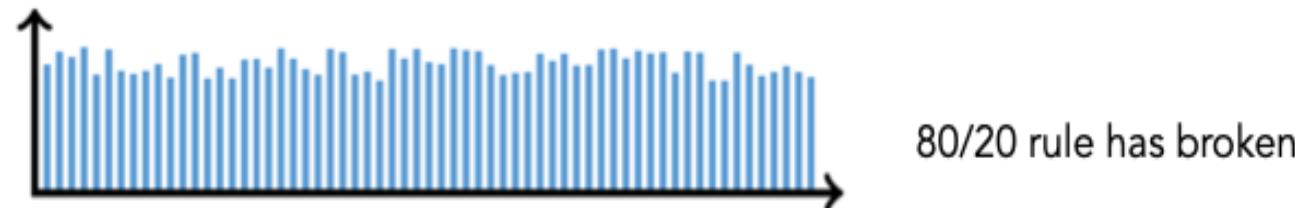
# Thrashing

## Page replacement algorithms avoid **thrashing**

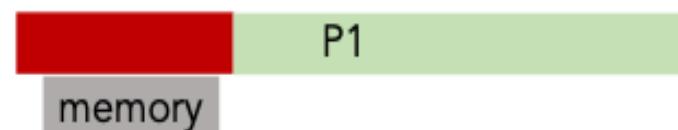
- When OS spent most of the time in paging data back and forth from disk
- Little time spent doing useful work (making progress)
- In this situation, the system is **overcommitted**
  - No idea which pages should be in memory to reduce faults
  - Ex: Running Windows95 with 4 MB of memory...

# Reason for Thrashing

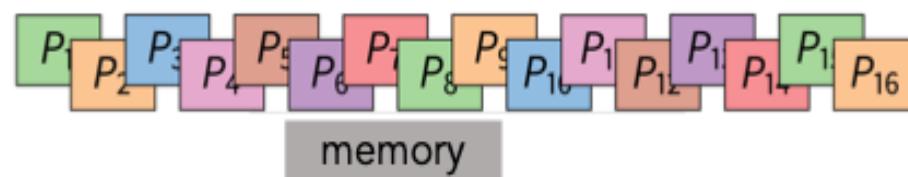
Access pattern has no temporal locality (past  $\approx$  future)



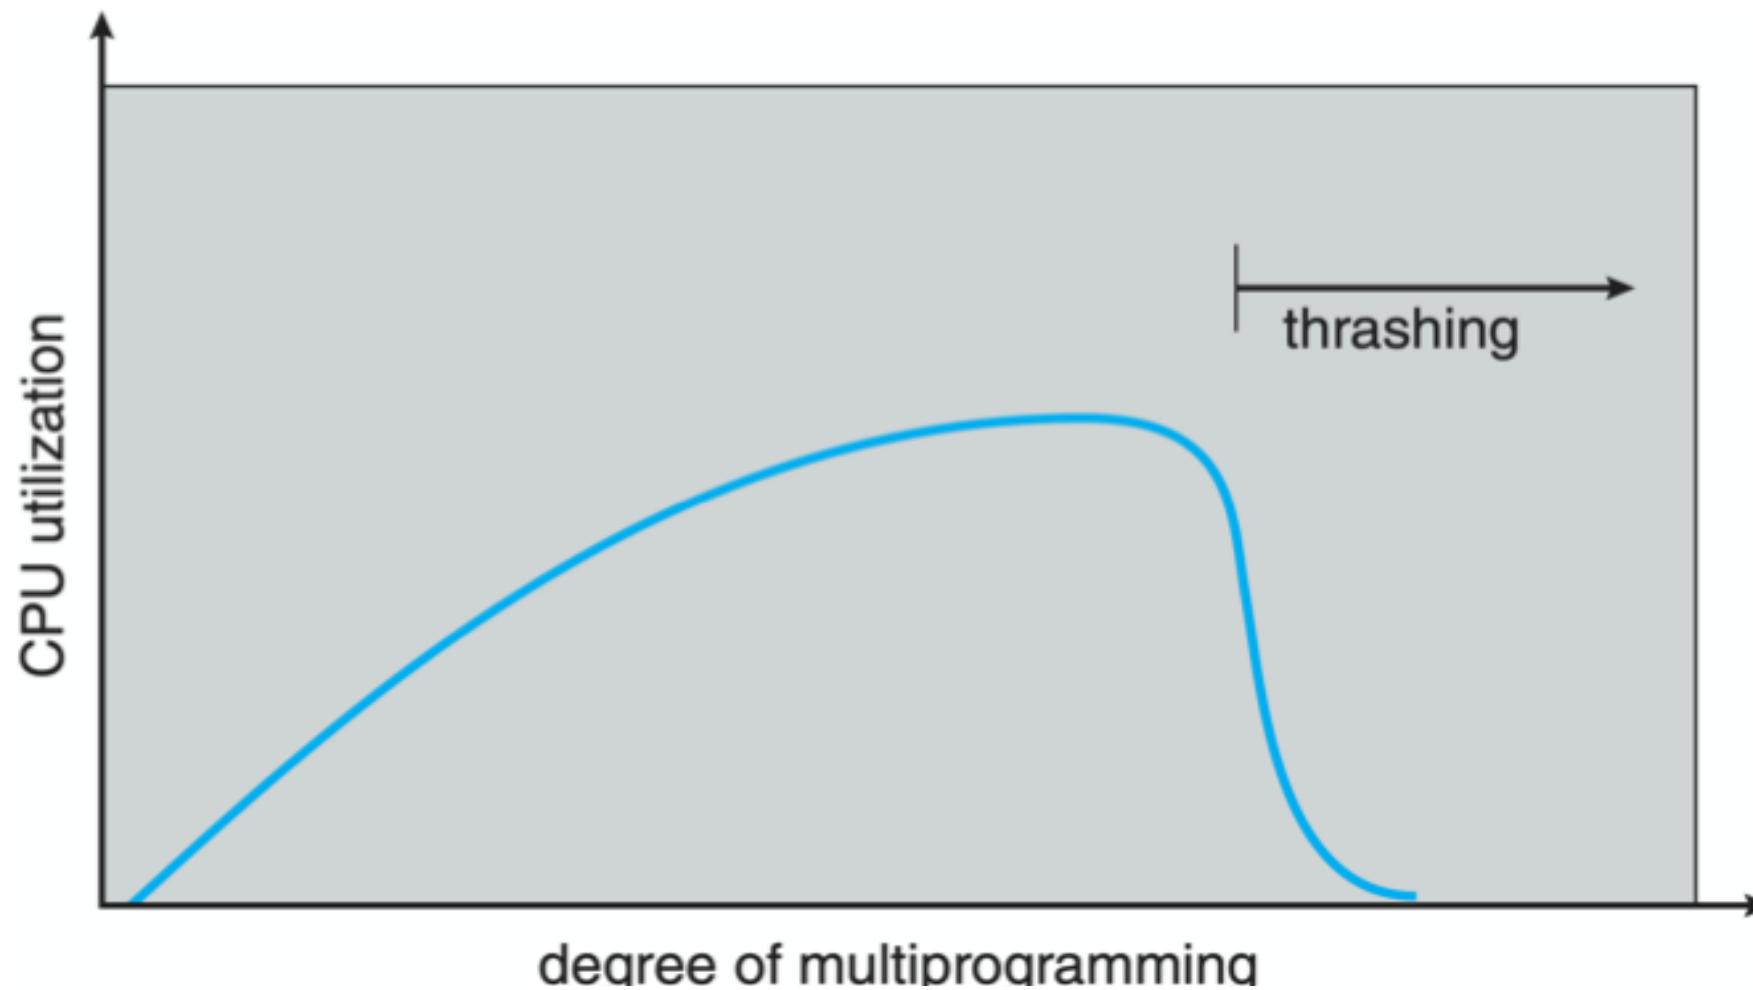
Hot memory does not fit in physical memory



Each process fits individually, but too many for system



# Thrashing & Multiprogramming



# Dealing with Thrashing

**Only run processes if memory requirements can be satisfied**

- Thrashing viewed from a caching perspective: given locality of reference, how big a cache does the process need?
- Or: how much memory does the process need in order to make reasonable progress (its working set)

**Swapping – write out all pages of a process**

**Buy more memory...**

# Summary

## Page replacement algorithms

- Belady's – optimal replacement (minimum # of faults)
- FIFO – replace page loaded furthest in past
- LRU – replace page referenced furthest in past
  - Approximate using PTE reference bit
- LRU Clock – replace page that is “old enough”
- Working Set – keep the set of pages in memory that has minimal fault rate
- Page Fault Frequency – grow/shrink page set as a function of fault rate

## Multiprogramming

- Should a process replace its own page, or that of another?

# **Next Time...**

**Read Chapter 14,17**