

# CE 440 Introduction to Operating System

Lecture 18: Fast File System  
Fall 2025

Prof. Yigong Hu



Slides courtesy of Manuel Egele, Ryan Huang and Baris Kasikci

# Midterm Result

Means: 39.78, Maximum: 60, std Dev: 9.23

Distribution:

> 60: 1

(50, 60]: 4

(45, 50]: 6

(40,45]: 14

(35,40]: 12

<35: 12

# Administrivia

**Midterm solution will not be directly posted online**

- The Rubric items are published in gradescope
- Can request to see the copy of sample solution in my office hour

**Lab 4 is out, this is optional**

# File Systems Examples

## BSD Fast File System (FFS)

- What were the problems with the original Unix FS?
- How did FFS solve these problems?

## Log-Structured File system (LFS) – next lecture

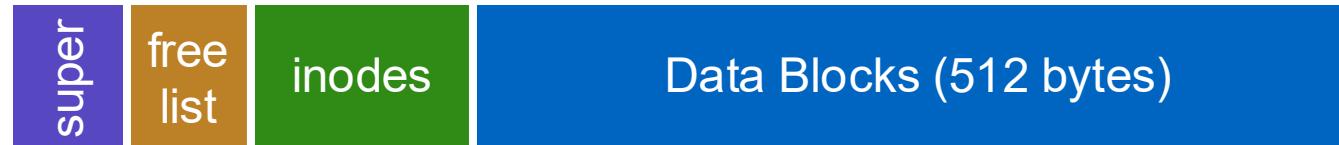
- What was the motivation of LFS?
- How did LFS work?

# Original Unix FS

From Bell Labs by Ken Thompson

Simple and elegant:

Unix disk layout



## Components

- Data blocks
- Inodes (directories represented as files)
- Free list
- Superblock. (specifies number of blks in FS, counts of max # of files, pointer to head of free list)

## Problem: slow

- Only gets **2% of disk maximum** (20Kb/sec) even for sequential disk transfers!

# Why So Slow?

## Problem 1: blocks too small (512 bytes)

- File index too large
- Require more indirect blocks
- Transfer rate low (get one block at time)

## Problem 2: unorganized freelist

- Consecutive file blocks not close together
  - Pay seek cost for even sequential access
- Aging: becomes fragmented over time

## Problem 3: poor locality

- inodes far from data blocks
- inodes for directory not close together
  - poor enumeration performance: e.g., “ls”, “grep foo \*.c”

# FFS: Fast File System

**Designed by a Berkeley research group for the BSD UNIX**

- A classic file systems paper to read: [[McKusic](#)]

**Approach:**

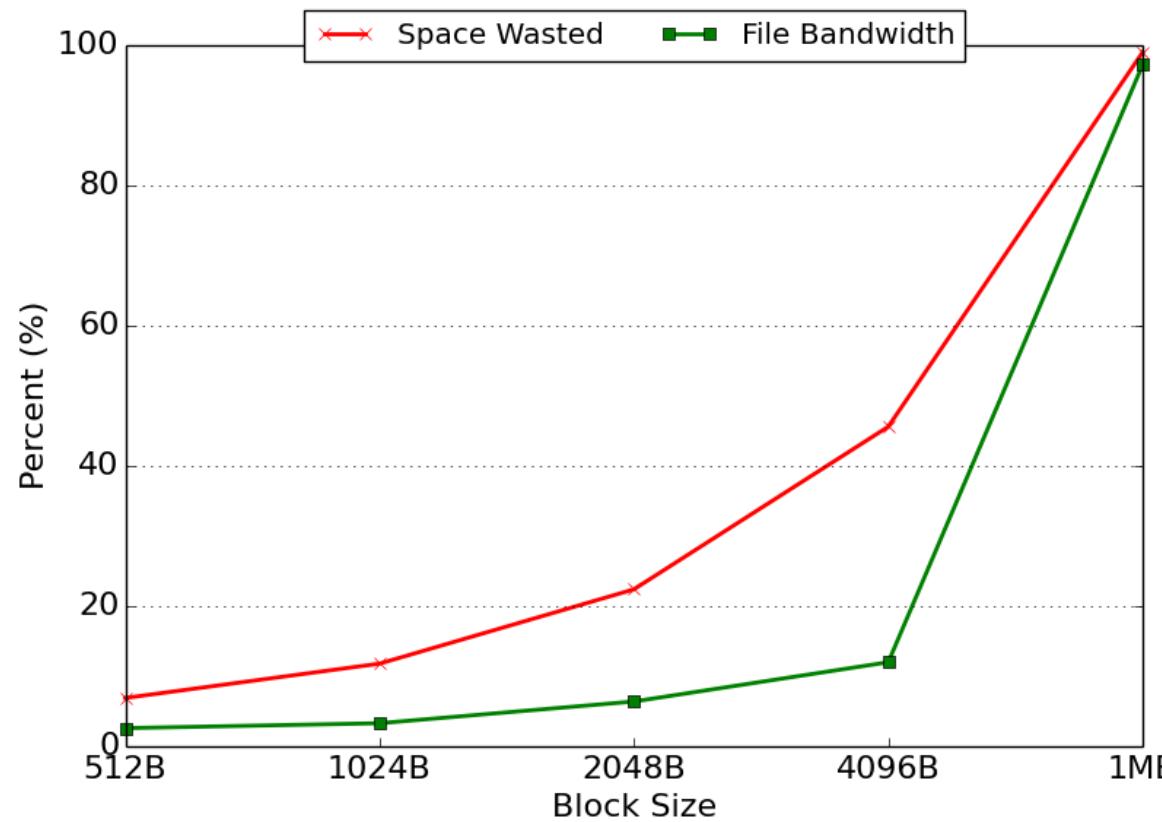
- [Measure](#) an state of the art systems
- Identify and understand the fundamental problems
  - The original FS treats disks like random-access memory!
- Get an idea and [build](#) a better systems

**Idea: design FS structures and allocation polices to be “disk aware”**

**Next: how FFS fixes the performance problems (to a degree)**

# Problem 1: Blocks Too Small

Measurement:



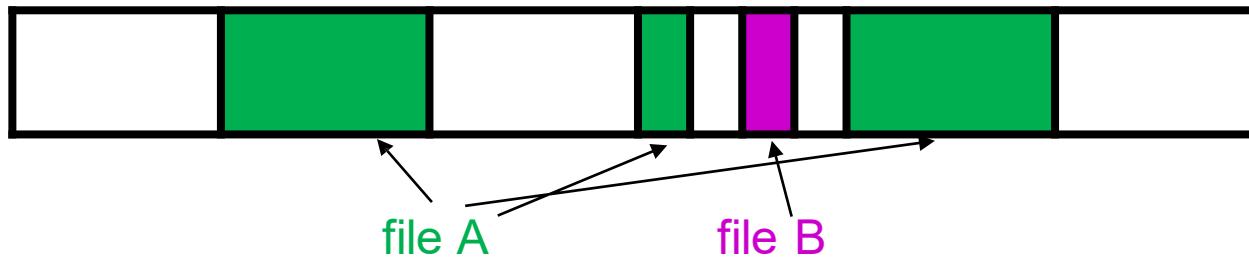
Bigger block increases bandwidth, but how to deal with wastage (“internal fragmentation”)?

- Use idea from malloc: split unused portion

# Solution: Fragments

## BSD FFS:

- Has large block size (4096B or 8192B)
- Allow large blocks to be chopped into small ones called “fragments”
- Ensure fragments only used for little files or ends of files



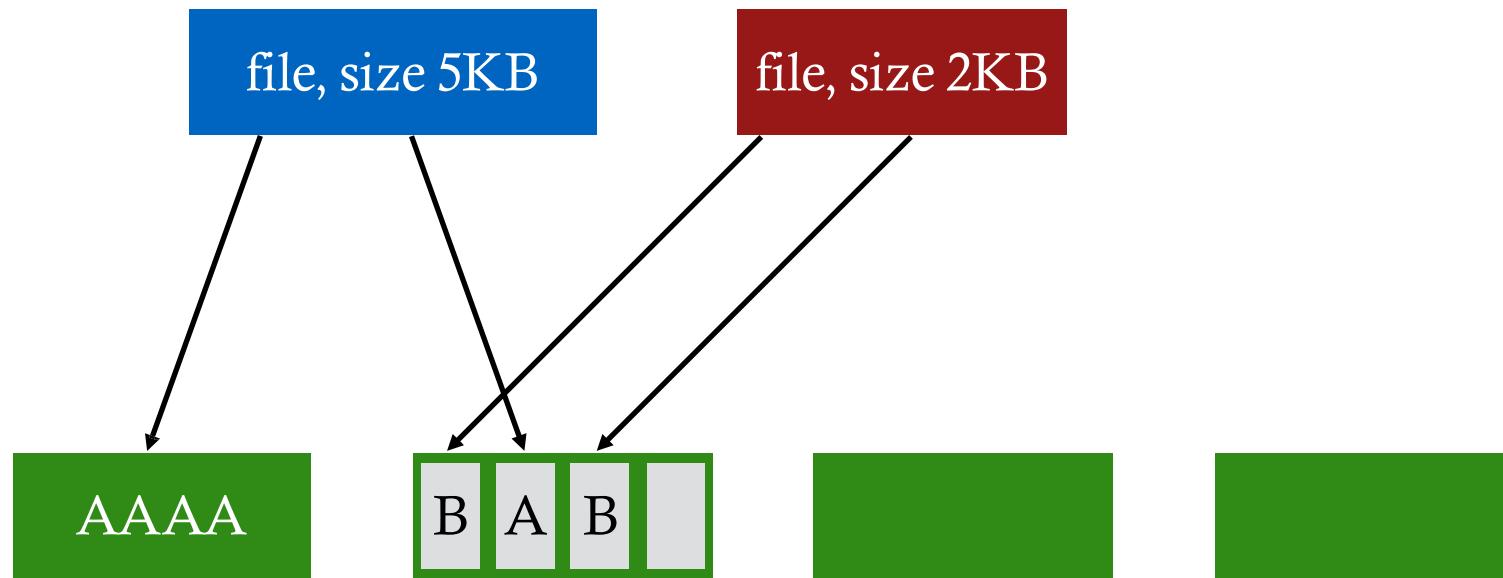
- Fragment size specified at the time that the file system is created
- Limit number of fragments per block to 2, 4, or 8

## Pros

- High transfer speed for larger files
- Low wasted space for small files or ends of files

# Fragment Example

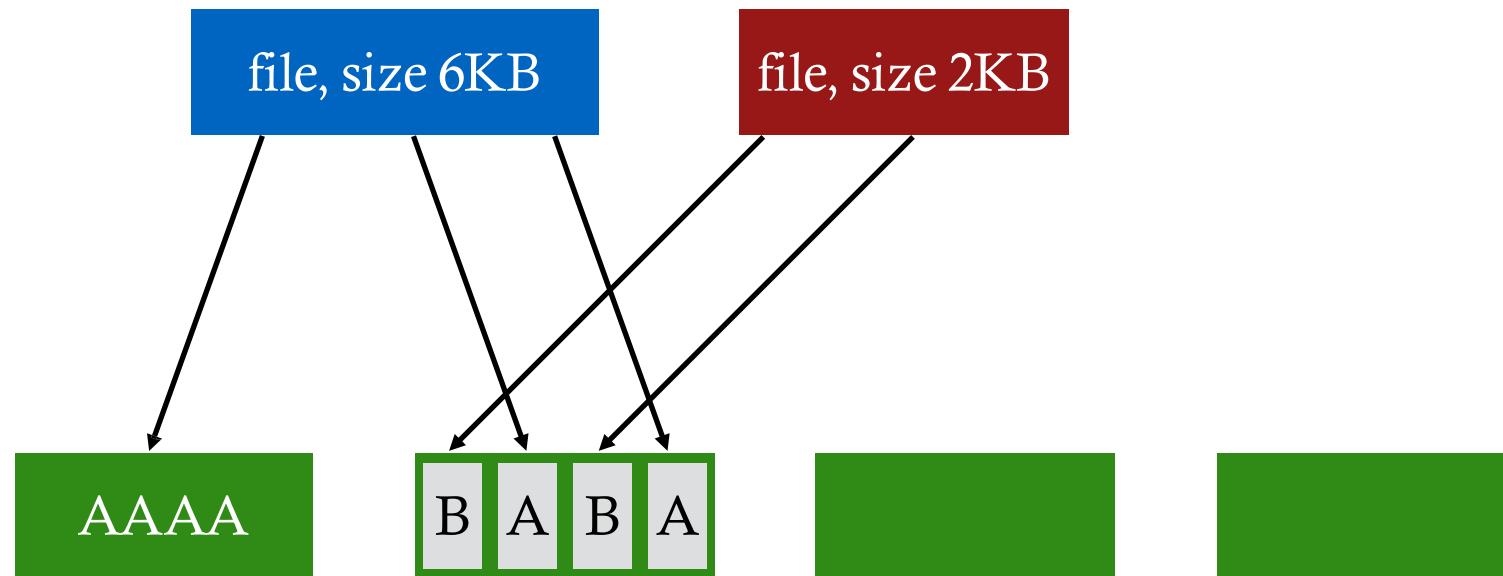
Block size: 4096 B  
Fragment size: 1024 B



# Fragment Example

```
write(fd1, "A"); // append A to first file
```

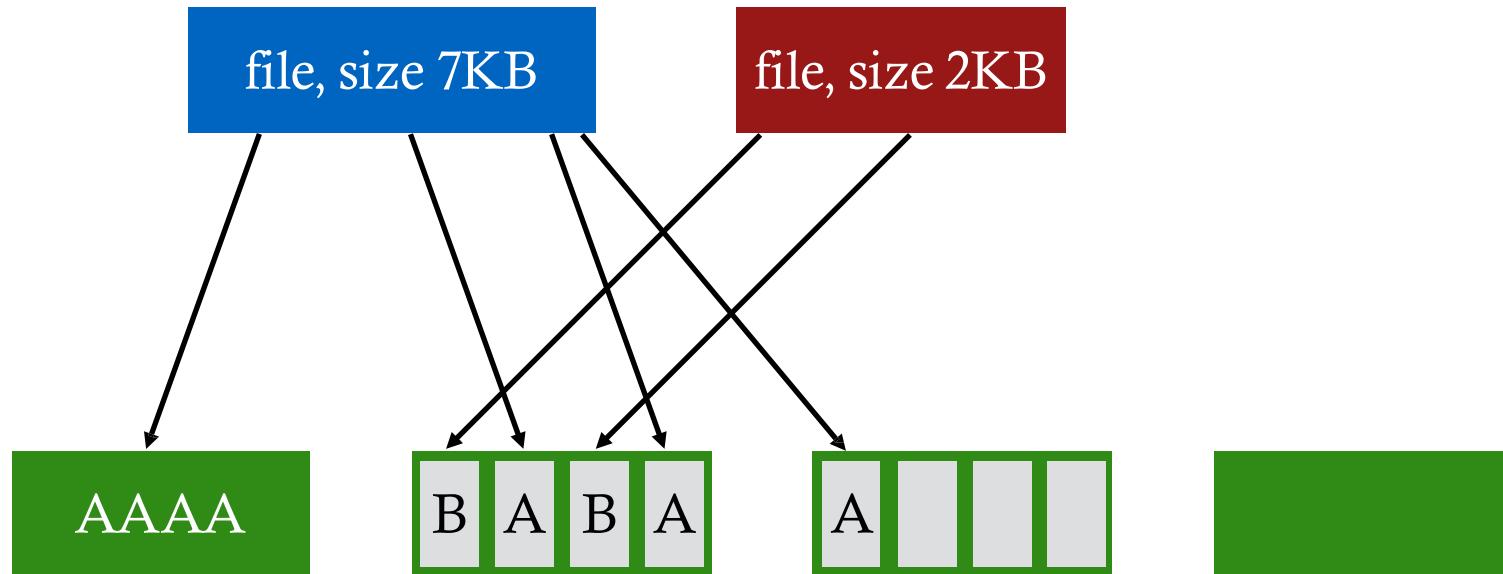
Block size: 4096 B  
Fragment size: 1024 B



# Fragment Example

```
write(fd1, "A"); // append A to first file  
write(fd1, "A");
```

Block size: 4096 B  
Fragment size: 1024 B



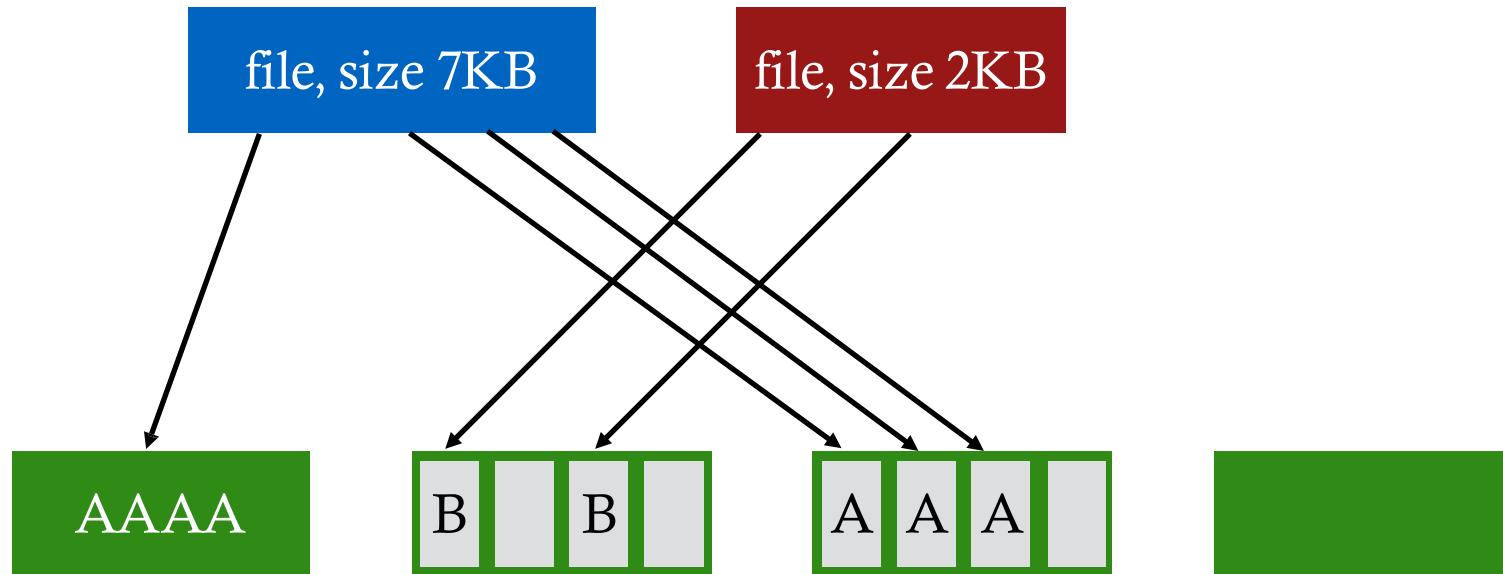
**Not allowed to use fragments across multiple blocks!**

What to do instead?

# Fragment Example

```
write(fd1, "A"); // append A to first file  
write(fd1, "A");
```

Block size: 4096 B  
Fragment size: 1024 B

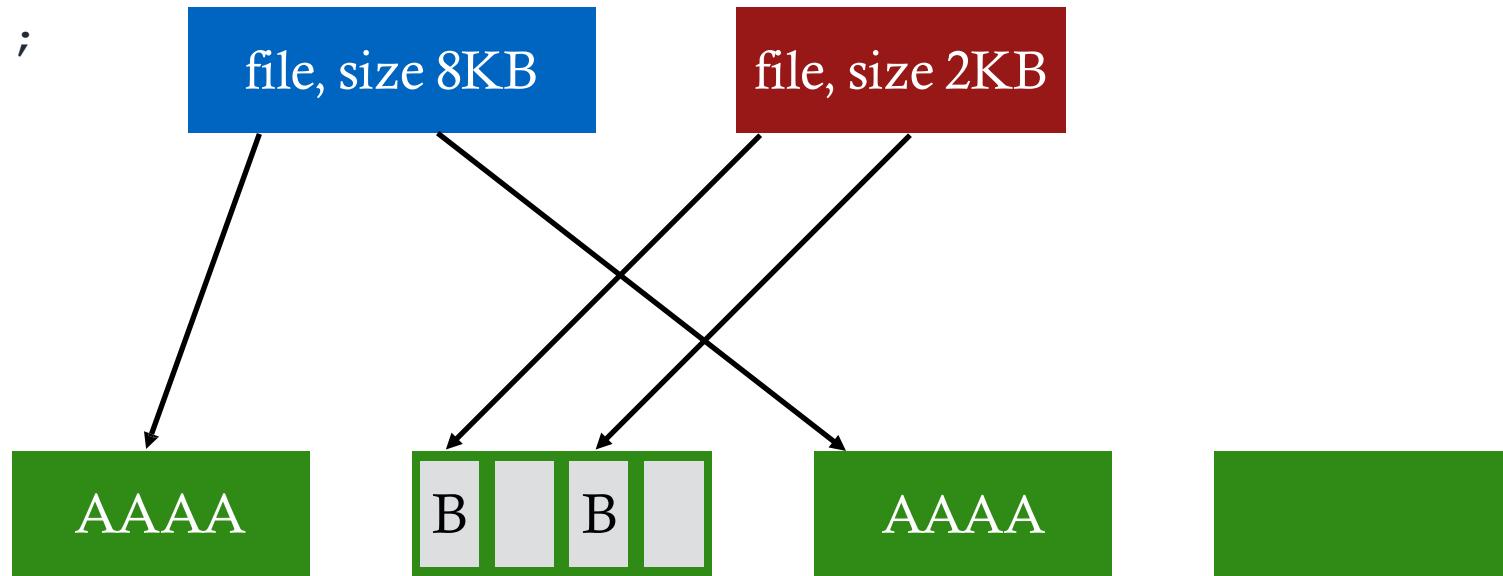


copy old fragments to new block  
new data use remaining fragments

# Fragment Example

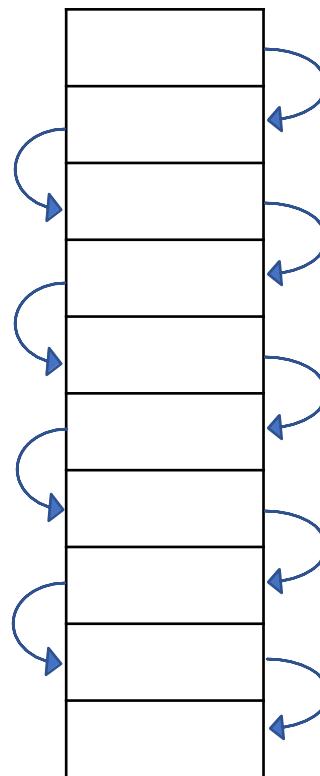
```
write(fd1, "A"); // append A to first file  
write(fd1, "A");  
write(fd1, "A");
```

Block size: 4096 B  
Fragment size: 1024 B

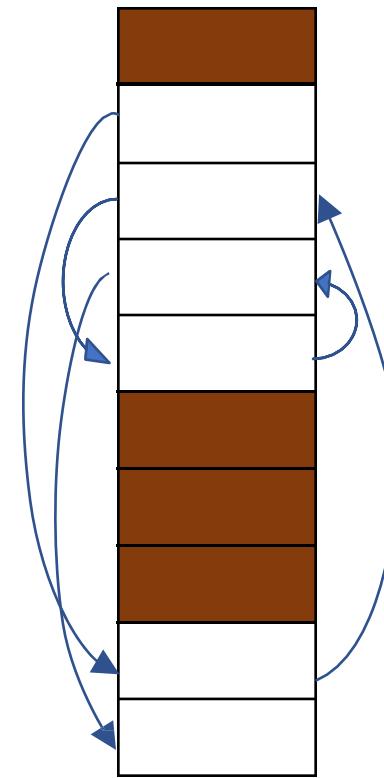


# Problem 2: Unorganized Freelist

Leads to random allocation of sequential file blocks overtime



Initial performance good



Get worse over time

## Measurement:

- New FS: 17.5% of disk bandwidth
- Few weeks old: 3% of disk bandwidth

# Fixing the Unorganized Freelist

## Periodical compact/defragment disk

- Cons: locks up disk bandwidth during operation

## Keep adjacent free blocks together on freelist

- Cons: costly to maintain

## FFS: bitmap of free blocks

- Each bit indicates whether block is free
  - E.g., 101010111111000001111111000101100
- Easier to find contiguous blocks
- Small, so usually keep entire thing in memory
- Time to find free blocks increases if fewer free blocks
- What about fragments in a block?

Bits in map	XXXX	XXOO	OOXX	0000
Fragment numbers	0-3	4-7	8-11	12-15
Block numbers	0	1	2	3

# Using a Bitmap

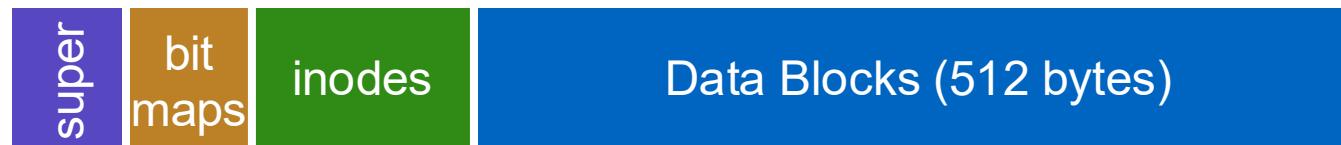
**Usually keep entire bitmap in memory:**

- 4G disk / 4K byte blocks. How big is map?

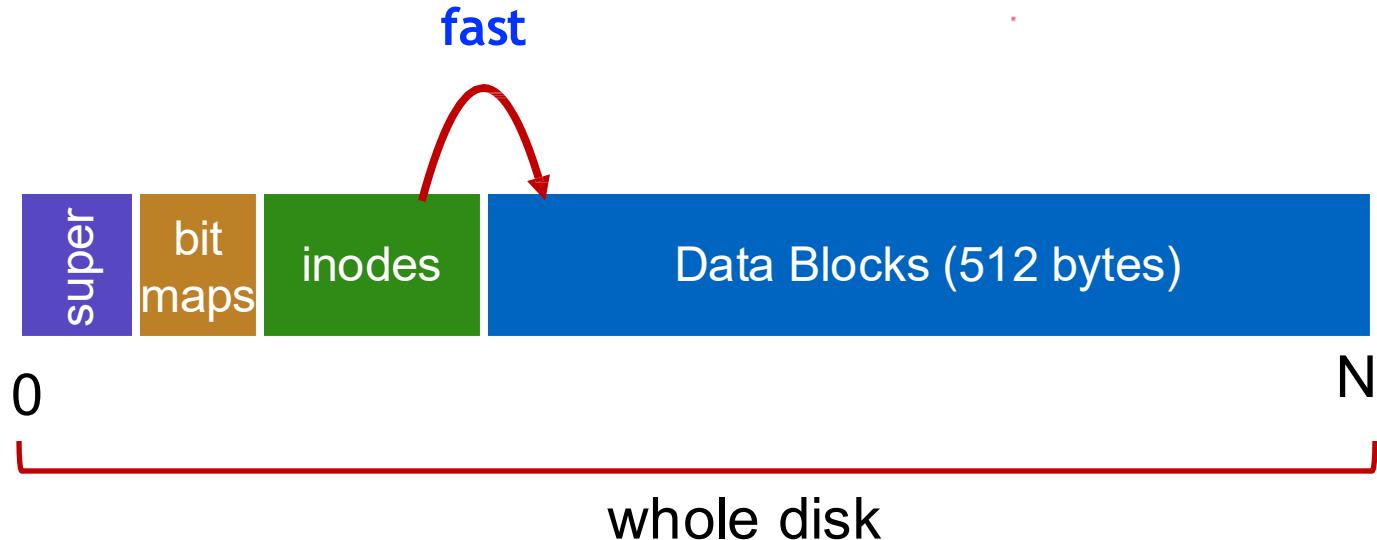
**Allocate block close to block x?**

- Check for blocks near  $bmap[x/32]$
- If disk almost empty, will likely find one near
- As disk becomes full, search becomes more expensive and less effective

**Trade space for time (search time, file access time)**

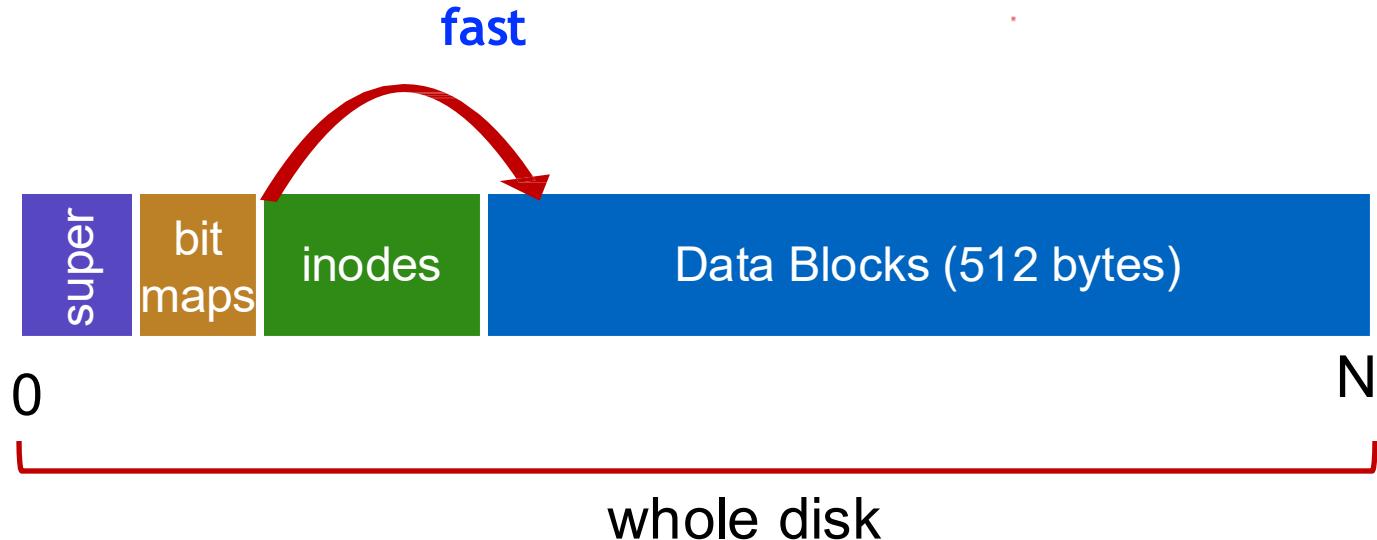


# Problem 3: Poor Locality



**How to keep inode close to data block?**

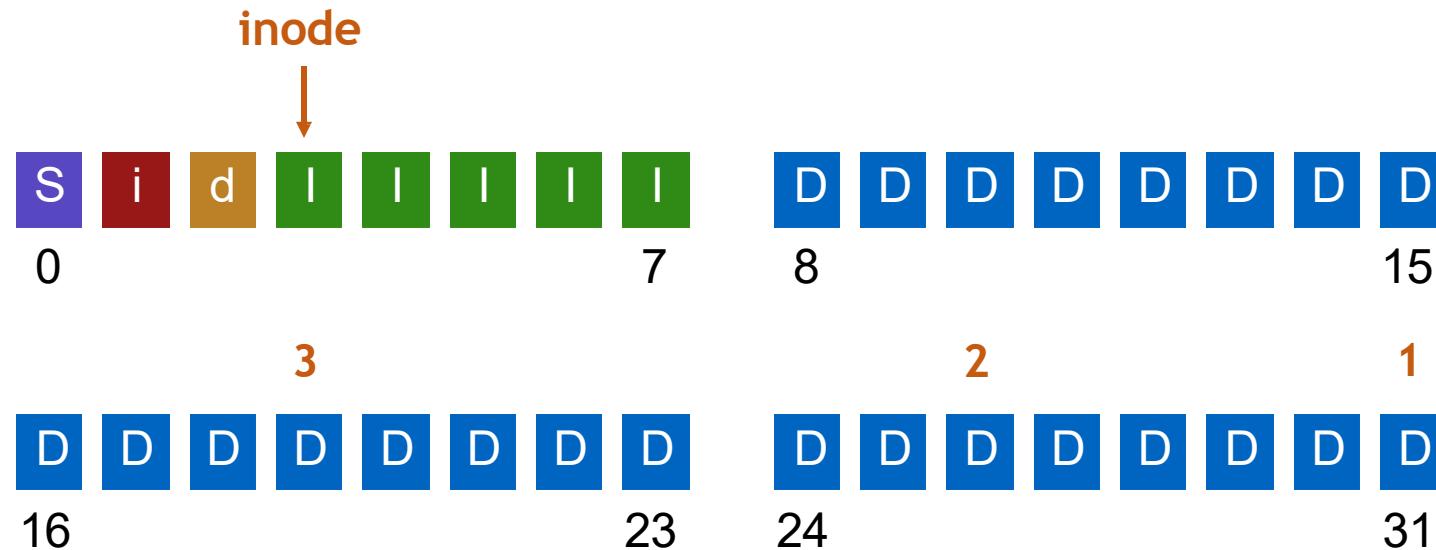
# Problem 3: Poor Locality



**How to keep inode close to data block?**

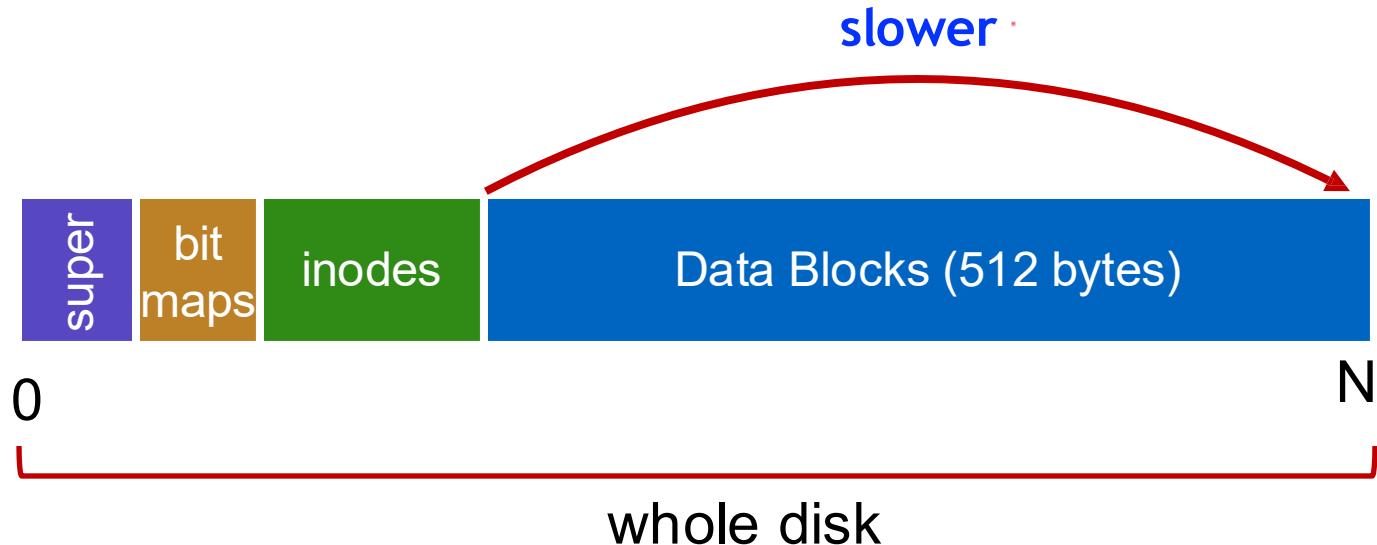
# Problem 3: Poor Locality

Example bad layout:



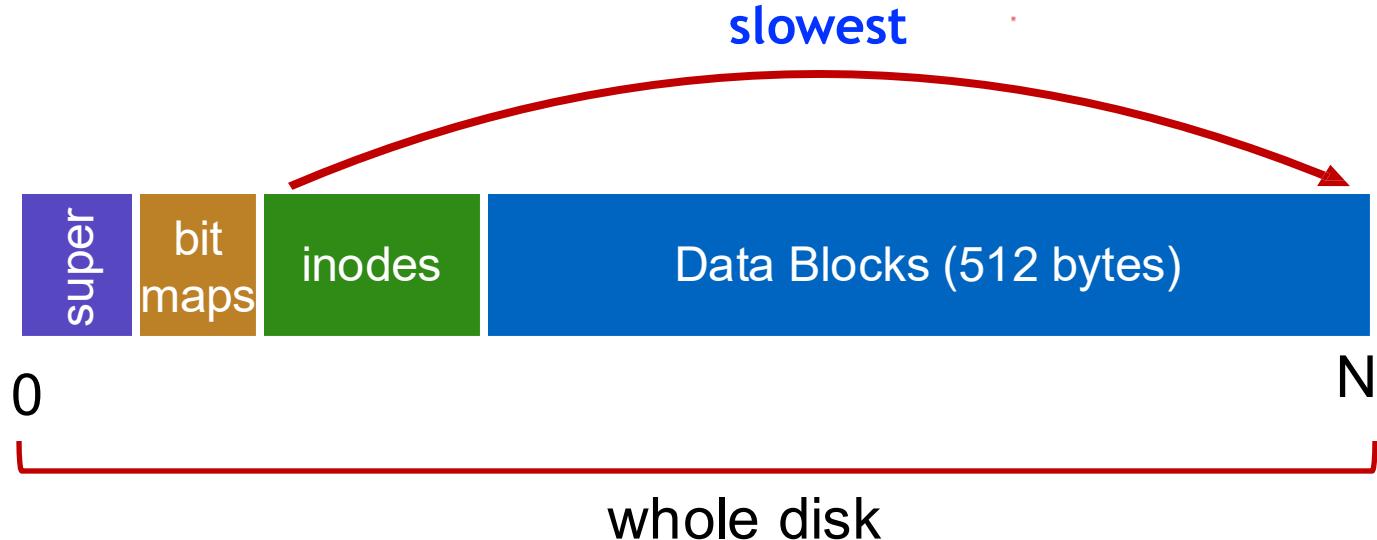
How to keep inode close to data block?

# Problem 3: Poor Locality



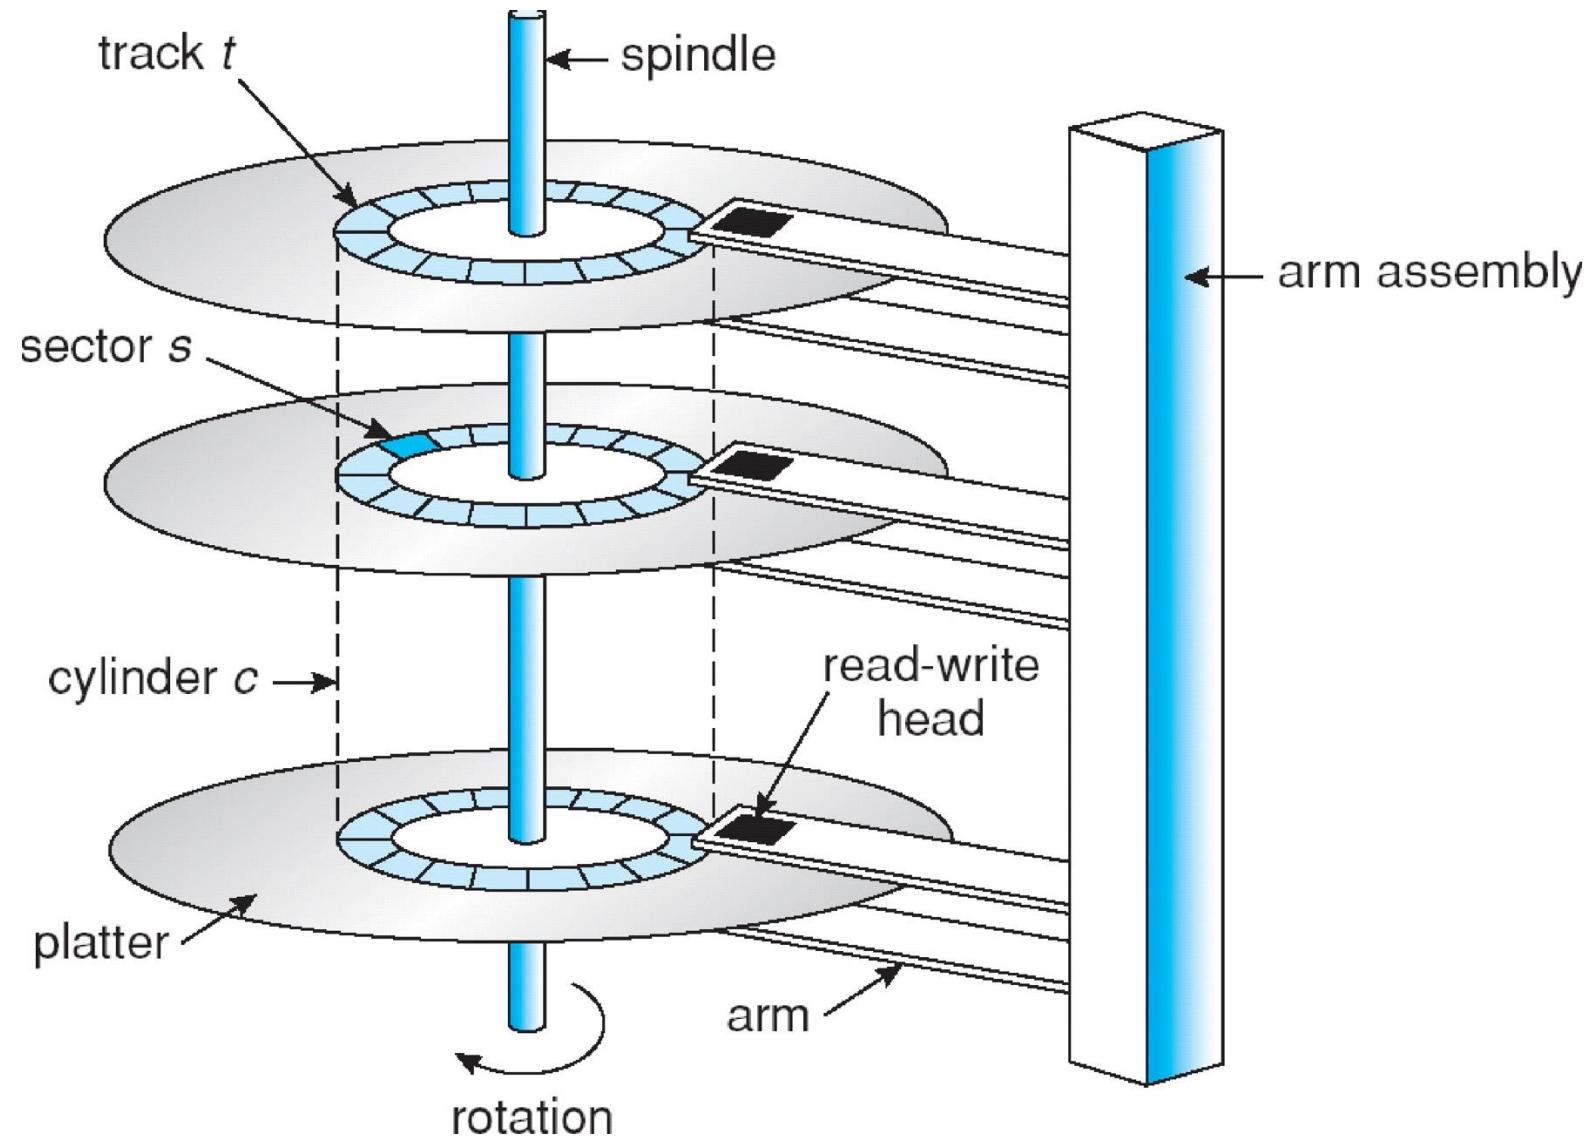
**How to keep inode close to data block?**

# Problem 3: Poor Locality



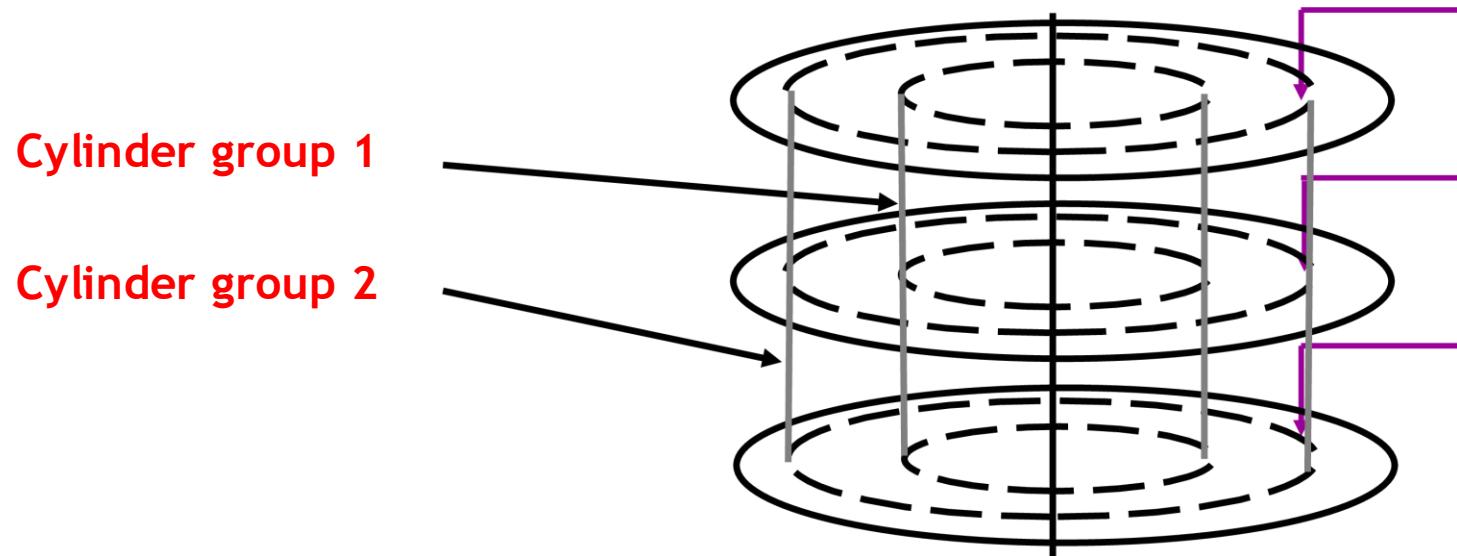
**How to keep inode close to data block?**

# Recap: Cylinders, Trackers, & Sectors



# FFS Solution: Cylinder Group

Group sets of consecutive cylinders into “cylinder groups”



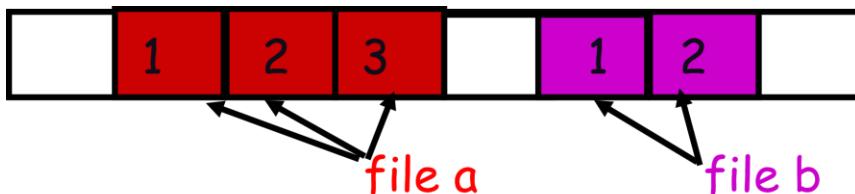
**Key: can access any block in a cylinder without performing a seek. Next fastest place is adjacent cylinder.**

- Tries to put everything related in same cylinder group
- Tries to put everything not related in different group

# Clustering in FFS

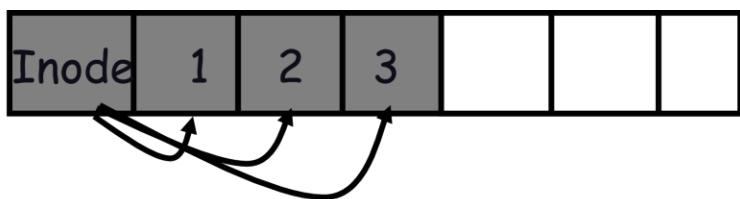
**Tries to put sequential blocks in adjacent sectors**

- (Access one block, probably access next)



**Tries to keep inode in same cylinder as file data:**

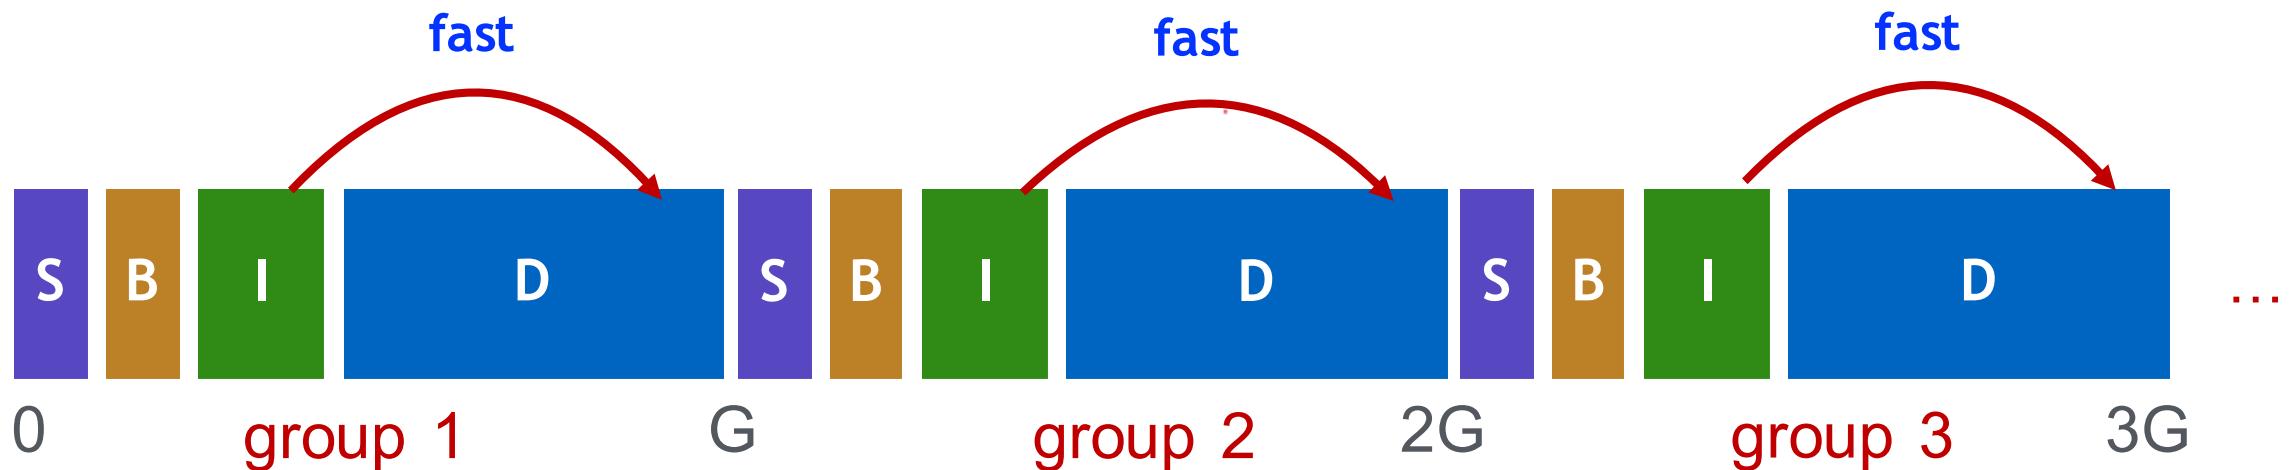
- (If you look at inode, most likely will look at data too)



**Tries to keep all inodes in a dir in same cylinder group**

- Access one name, frequently access many, e.g., “ls -l”

# What Does Disk Layout Look Like Now?



**How to keep inode close to data block?**

- Answer: Use groups across disks
- Strategy: allocate inodes and data blocks in same group
- Each cylinder group basically a mini-Unix file system

**Is it useful to have multiple super blocks?**

- Yes, if some (but not all) fail

# FFS Results

## Performance improvements:

- Able to get 20-40% of disk bandwidth for large files
- 10-20x original Unix file system!
- Stable over FS lifetime
- Better small file performance (why?)

## Other enhancements

- Long file names
- Parameterization
- Free space reserve (10%) that only admin can allocate blocks from

# **Next Time...**

Read Chapter 43