

CE 440 Introduction to Operating System

Midterm Review
Fall 2025

Prof. Yigong Hu



Slides courtesy of Manuel Egele, Ryan Huang and Baris Kasikci

Midterm

- **October 22th Tuesday 4:30-6:00 pm at classroom**
- **Covers material in lecture1-9**
- **Based upon lecture material, homework, and project**
 - Make sure you do the homework to practice
- **One 8.5''x11" double-sided sheet of notes**
- **Obligatory: do not cheat**
 - Everyone must put all of their belongings in a bag and place the bag at the front of the room.
 - Do not copy from your neighbors
 - No one involved will be happy, particularly the teaching staff

Exam Format

- **7 questions (75 points + 5 bonus point)**
 - 1 true/false question
 - 4 short answer questions
 - 1 coding question
 - 1 lab related question

Lecture Overview

Architecture support

Process + Thread

Scheduling

Synchronization

Deadlock

Arch Support for OSes

Types of architecture support

- Manipulating privileged machine state
- Generating and handling events

Privileged Instructions

What are privileged instructions?

- Who gets to execute them?
- How does the CPU know whether they can be executed?
- Difference between user and kernel mode

Why do they need to be privileged?

What do they manipulate?

- Protected control registers
- Memory management
- I/O devices

Events

Events

	Unexpected	Deliberate
Exceptions (sync)	fault	syscall trap
Interrupts (async)	interrupt	Software interrupt

What are faults, and how are they handled?

What are system calls, and how are they handled?

What are interrupts, and how are they handled?

- How do I/O devices use interrupts?

What is the difference between exceptions and interrupts?

Lecture Overview

Architecture support

Process + Thread

Scheduling

Synchronization

Deadlock

Processes

What is a process?

What resource does it virtualize?

What is the difference between a process and a program?

What is contained in a process?

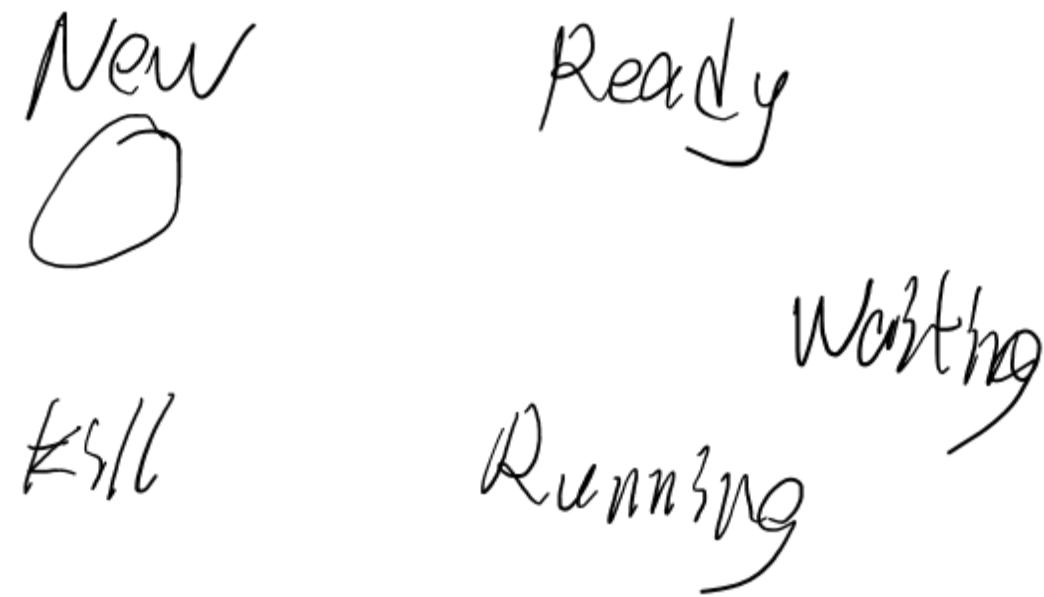
Processes Data Structures

Process Control Blocks (PCBs)

- What information does it contain?
- How is it used in a context switch?

State queues

- What are process states?
- What is the process state graph?
- When does a process change state?
- How does the OS use queues to keep track of processes?



Processes Manipulation

What does CreateProcess on NT do?

What does fork() on Unix do?

- What does it mean for it to “return twice”?

What does exec() on Unix do?

- How is it different from fork?

How are fork and exec used to implement shells?

Why fork()?

Threads

What is a thread?

- What is the difference between a thread and a process?
- How are they related?

Why are threads useful?

What is the difference between user-level and kernel-level threads?

- What are the advantages/disadvantages of one over another?

Threads Implementation

How are threads managed by the run-time system?

- Thread control blocks, thread queues
- How is this different from process management?

What operations do threads support?

- create, yield, sleep, etc.
- What does thread yield do?

What is a context switch?

What is the difference between non-preemptive scheduling and preemptive thread scheduling?

- Voluntary and involuntary context switches

Lecture Overview

Architecture support

Process + Thread

Scheduling

Synchronization

Deadlock

Scheduling

What kinds of scheduling is there?

- Long-term scheduling
- Short-term scheduling

Components

- Scheduler (dispatcher)

When does scheduling happen?

- Job changes state (e.g., waiting to running)
- Interrupt, exception
- Job creation, termination

Scheduling Goals

Goals

- Maximize CPU utilization
- Maximize job throughput
- Minimize turnaround time
- Minimize waiting time
- Minimize response time

What is the goal of a batch system?

What is the goal of an interactive system?

Starvation

Starvation

- Indefinite denial of a resource (CPU, lock)

Causes

- Side effect of scheduling
- Side effect of synchronization

Operating systems try to prevent starvation

Scheduling Algorithm

What are the properties, advantages and disadvantages of the following scheduling algorithms?

- First Come First Serve (FCFS)/First In First Out (FIFO)
- Shortest Job First (SJF)

Preemptive: Shortest-Remaining-Time-First (SRTF)

- Priority
- Round Robin
- Multilevel feedback queues

What scheduling algorithm does Unix use? Why?

Lecture Overview

Architecture support

Process + Thread

Scheduling

Synchronization

Deadlock

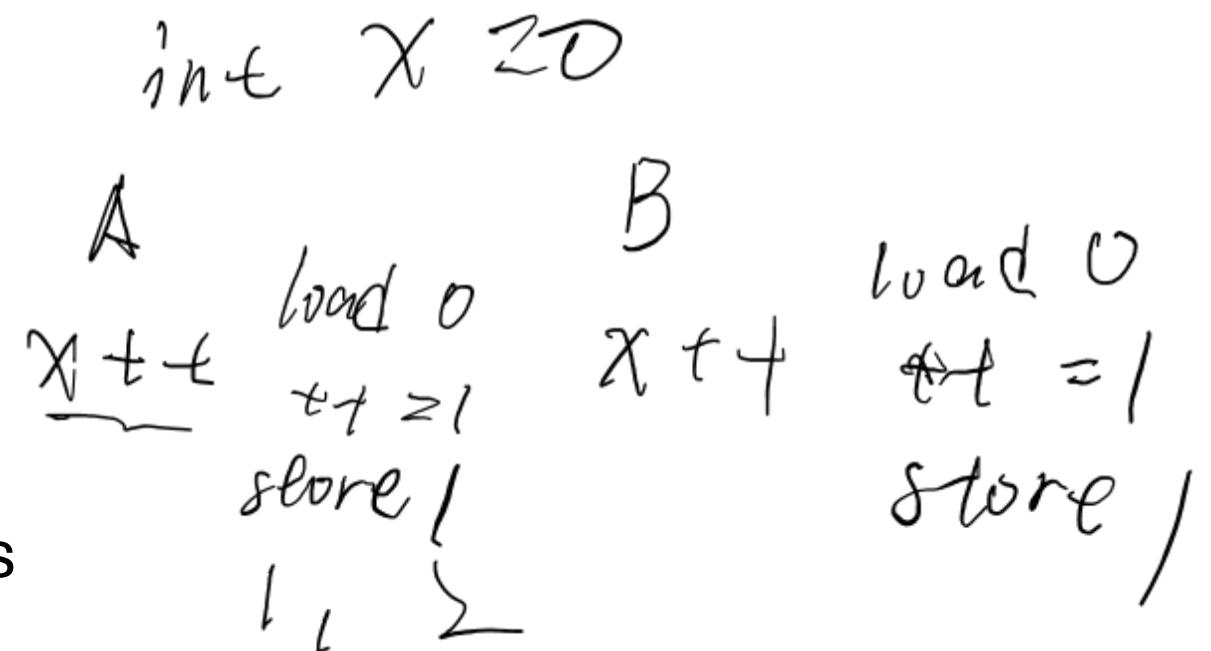
Synchronization

Why do we need synchronization?

- Coordinate access to shared data structures
- Coordinate thread/process execution

What can happen to shared data structures if synchronization is not used?

- Race condition
- Corruption
- Bank account example



When are resources shared?

- Global variables, static objects
- Heap objects

Concurrent Programs

```
Monitor Bounded_buffer {
    Resource buffer[N];
    // Variables for indexing buffer
    // monitor invariant involves these vars
    Condition not_full; // space in buffer
    Condition not_empty; // value in buffer

    void put_resource (Resource R) {
        while (buffer array is full)
            wait(not_full);
        Add R to buffer array;
        signal(not_empty);
    }
}
```

```
Resource get_resource() {
    while (buffer array is empty)
        wait(not_empty);
    Get resource R from buffer array;
    signal(not_full);
    return R;
}

} // end monitor
```

Our goal is to write concurrent programs...

Concurrent Programs

Need mutual exclusion for critical sections

```
Resource get_resource() {  
    while (buffer array is empty)  
        wait(not_empty);  
    Get resource R from buffer array;  
    signal(not_full);  
    return R;  
}  
} // end monitor
```

Need mechanisms for coordinating threads

Mutual Exclusion

Need mutual exclusion for critical sections

```
lock.acquire();  
...  
lock.release();
```

Interrupts enabled, other threads can run (just not in this critical section)

Mutual Exclusion

```
Void acquire () {  
    // Disable interrupts  
    // Enable interrupts  
}
```

```
lock.acquire();  
  
...  
  
lock.release();
```

Also need mutual exclusion; disable
interrupts, or use spinlocks with special
hardware instructions

Mutual Exclusion

What is mutual exclusion?

What is a critical section?

- What guarantees do critical sections provide?
- What are the requirements of critical sections?
 - Mutual exclusion (safety)
 - Progress (liveness)
 - Bounded waiting (no starvation: liveness)
 - Performance

How does mutual exclusion relate to critical sections?

What are the mechanisms for building critical sections?

- Locks, semaphores, monitors, condition variables

Locks

What does Acquire do?

What does Release do?

What does it mean for Acquire/Release to be atomic?

How can locks be implemented?

- Spinlocks
- Disable/enable interrupts
- Blocking

How does test-and-set work?

- What kind of lock does it implement?

What are the limitations of using spinlocks, interrupts?

- Inefficient, interrupts turned off too long

Semaphores

What is a semaphore?

- What does Wait/P/Decrement do?
- What does Signal/V/Increment do?
- How does a semaphore differ from a lock?
- What is the difference between a binary semaphore and a counting semaphore?

When do threads block on semaphores?

When are they woken up again?

Using semaphores to solve synchronization problems

- Readers/Writers problem
- Bounded Buffers problem

Monitors

What is a monitor?

- Shared data
- Procedures
- Synchronization

In what way does a monitor provide mutual exclusion?

- To what extent is it provided?

How does a monitor differ from a semaphore?

How does a monitor differ from a lock?

What kind of support do monitors require?

- Language, run-time support

Condition Variables

What is a condition variable used for?

- Coordinating the execution of threads
- Not mutual exclusion

Operations

- What are the semantics of Wait?
- What are the semantics of Signal?
- What are the semantics of Broadcast?

How are condition variables different from semaphores?

Implementing Monitors

What does the implementation of a monitor look like?

- Shared data
- Procedures
- A lock for mutual exclusion to procedures (w/ a queue)
- Queues for the condition variables

What is the difference between Hoare and Mesa monitors?

- Semantics of signal (whether the woken up waiter gets to run immediately or not)
- What are their tradeoffs?
- What does Java provide?

Locks and Condition Variables

Condition variables are also used without monitors in conjunction with locks

A monitor ≈ a module whose state includes a C/V and a lock

Why must cond_wait both release mutex_t & sleep?

Synchronization

Use synchronization primitives (locks, semaphores, monitor, condition variables, etc.) to solve synchronization problems

Lecture Overview

Architecture support

Process + Thread

Scheduling

Synchronization

Deadlock

Deadlock

Deadlock happens when processes are waiting on each other and cannot make progress

What are the conditions for deadlock?

- Mutual exclusion
- Hold and wait
- No preemption
- Circular wait

How to visualize, represent abstractly?

- Resource allocation graph (RAG)
- Waits for graph (WFG)

Deadlock Approaches

Dealing with deadlock

- Ignore it
- Prevent it (prevent one of the four conditions)
- Avoid it (have tight control over resource allocation)
- Detect and recover from it

What is the Banker's algorithm?

- Which of the four approaches above does it implement?

Race Conditions

```
int x = 0;
int i, j;
void AddToX() {
    for (i = 0; i < 100; i++) x++;
}

void SubFromX() {
    for (j = 0; j < 100; j++) x--;
}
```

What is the range of possible values for x? Why?