

EC440 Pintos Project Lab2

Overview

William Wang (xwill@bu.edu)

Administrivia

- **Lab 2 deadline: Friday 11/07 11:59 pm**
 - Estimated time: **50~60** hours per group
- **Submission**
 - create a “lab2-handin” branch

Outline

→ **User Programs In Pintos**


- An Overview of Project 2 Requirements
- Getting start
- Tips

Overview

- **Project 2: Userprog**
 - Allow **user programs** to run on top of pintos
- **Lab 2 requires good **understanding** of**
 - How user programs run in general
 - Distinctions between user and kernel virtual memory
 - System call infrastructure
 - File system interface

User Process & Syscalls

- **Syscalls provide the interface between user process and OS**



How to Run A User Program In Linux

- What happens when a user runs (in the shell)

- cp -r ~/foo ~/foo1

- Shell parses user input

- - argc = 4, argv = "cp", "-r", "~/foo", "~/foo1"

- shell calls fork() and execve("cp", argv, env)

Load user program

- cp uses file system interface to copy files

User program calls syscall

- cp may print messages to stdout

- cp exits

User Programs In Pintos

- **Pintos implements a basic program loader**
 - Parse and load ELF executables
 - Start executables as a user process with one thread
- **But this system has problems (fixed by you in this lab)**
 - User processes crash immediately :(
 - System calls only print “system call!”

How Does Pintos Start A User Program?

- In **threads/init.c**
 - `pintos_init() -> run_actions() --> run_task(argv)`
 - `run_task() --> process_wait(process_execute(task))`


```
284     /* Runs the task specified in ARGV[1]. */
285     static void
286     run_task (char **argv)
287     {
288         const char *task = argv[1];
289
290         printf ("Executing '%s'\n", task);
291 #ifdef USERPROG
292         process_wait (process_execute (task));
293 #else
294         run_test (task);
295 #endif
296         printf ("Execution of '%s' complete.\n", task);
297     }
```

Execute user
program & wait

Important Process Functions In Pintos

- **The `process_execute()` start a process:**
 - creates thread running `start_process()`
 - `start_process()` thread loads executable file
 - sets up user **virtual memory** (stack, data, code)
 - starts executing user process (jump to the start)
- **`process_wait()` waits for executable to finish**
- **`process_exit()` frees resources of program**

Pintos Program Loading Flowchart



User Program Startup

- After Pintos loads executables, it jumps to user process

```
static void start_process(void *file_name) {
    struct intr_frame if_ = /* initialize */;
    asm volatile ("movl %0, %esp; jmp intr_exit" : : "g" (&if_) : "memory");
    NOT_REACHED ();
}
```

- _start() in lib/user/entry.c is entry point of user programs

```
void _start (int argc, char *argv[]) {
    exit (main (argc, argv))
}
```


- Kernel must pass process **start arguments** on user stack

How Does Pintos Handle Syscall?

- **Pintos uses int 0x30 for system calls**
- **Pintos has code for dispatching syscalls from user programs**
 - i.e. user processes will push 1) syscall number and 2) arguments onto the stack and execute int 0x30
- **In the kernel, calling syscall_handler() in userprog/syscall.c**

```
static void syscall_handler (struct intr_frame *f) {  
    printf ("system call!\n");  
    thread_exit ();  
}
```

Syscall Flowchart (exit)




Outline

- User Programs In Pintos
- An Overview of Project 2 
- Getting Start
- Tips

Project 2 Requirements

- **In Project 2, you need to implement:**
 - Process exit messages
 - Argument Passing
 - System calls (Major)
 - Safe memory access
 - Denying write to in-use executable files (*Extra Credit*)


Lab 2 Structure



Argument Passing

- A user process starts at "int main(int argc, char** argv)"
- In preparation to start a user process, the kernel must
 - parse the command
"/bin/ls -l foo bar" => "/bin/ls", "-l", "foo", "bar"
 - push function's arguments onto the stack
- Implement the string parsing however you like in load()
 - strtok_r(...) in lib/string.c is helpful

User vs Kernel Virtual Memory



Setting Up User Stack

- **userprog/process.c**

```
/* Create a minimal stack by mapping a zeroed page at the top of user virtual memory.  
*/  
static bool setup_stack (void **esp) {  
    uint8_t *kpage;  
    bool success = false;  
  
    kpage = palloc_get_page (PAL_USER | PAL_ZERO);  
    if (kpage != NULL) {  
        success = install_page (((uint8_t *) PHYS_BASE) - PGSIZE, kpage, true);  
        if (success) *esp = PHYS_BASE;  
        else palloc_free_page (kpage);  
    }  
    return success;  
}
```

Get one page from page pool,
return kernel virtual address

Map the **user** virtual
address to the page

Set the stack point

You need to place argc and
*argv on the initial stack, since
they are parameters to main()

Argument Passing (Stack)

- Push the words onto the stack
- Word-align
- Push a null pointer sentinel
- Push the address of each word in right-to-left order
- Push argv and argc
- Push 0 as a fake return address

Address	Name	Data	Type
0xbfffffc	argv[3][...]	"bar\0"	char[4]
0xbfffff8	argv[2][...]	"foo\0"	char[4]
0xbfffff5	argv[1][...]	"-1\0"	char[3]
0xbfffffed	argv[0][...]	"/bin/ls\0"	char[8]
0xbfffffec	word-align	0	uint8_t
0xbfffffe8	argv[4]	0	char *
0xbfffffe4	argv[3]	0xbfffffc	char *
0xbfffffe0	argv[2]	0xbfffff8	char *
0xbfffffdc	argv[1]	0xbfffff5	char *
0xbfffffd8	argv[0]	0xbfffffed	char *
0xbfffffd4	argv	0xbfffffd8	char **
0xbfffffd0	argc	4	int
0xbfffffc	return address	0	void (*) ()

Design Tips For Argument Passing

- **Implement user stack push function for argument passing**
 - lib/string.c is helpful
- **Distinguish user virtual address and kernel virtual address when you are coding**
- **hex_dump() function is useful for seeing the layout of stack**
 - `void hex_dump (uintptr_t ofs, const void *buf_, size_t size, ...)`
 - Dumps the SIZE bytes in BUF to the console
 - `ofs` is the starting address of buf

Safe Memory Access

- Kernel may access memory through user-provided pointers
 - E.g. read(), write()
- This is dangerous!
 - null pointers
 - pointers to unmapped virtual addresses
 - pointers to kernel addresses
- In lab 2, you need to support reading from and writing to user memory for system calls that only access valid address

Two Approaches To Solving Memory Access

- **Approach #1 (simplest): verify every user pointer **before** dereferencing.**
 - Check in user address space (< PHYS_BASE)
 - Check mapped (using pagedir_get_page() in userprog/pagedir.c)
- **Approach #2: Modify **page fault handler** in exception.c**
 - Check in user address space (< PHYS_BASE)
 - Dereference. Invalid pointers will trigger page faults
 - More convenient for lab 3

Two Approaches To Solving Memory Access

```
/* Reads a byte at user virtual address UADDR.
```

```
UADDR must be below PHYS_BASE.
```

```
Returns the byte value if successful, -1 if a segfault  
occurred. */
```

```
static int  
get_user (const uint8_t *uaddr)  
{  
    int result;  
    asm ("movl $1f, %0; movzbl %1, %0; 1:"  
        : "=a" (result) : "m" (*uaddr));  
    return result;  
}
```

```
/* Writes BYTE to user address UDST.
```

```
UDST must be below PHYS_BASE.
```

```
Returns true if successful, false if a segfault occurred. */
```

```
static bool  
put_user (uint8_t *udst, uint8_t byte)  
{  
    int error_code;  
    asm ("movl $1f, %0; movb %b2, %1; 1:"  
        : "=a" (error_code), "=m" (*udst) : "q" (byte));  
    return error_code != -1;  
}
```

<https://stackoverflow.com/questions/14922022/need-to-figure-out-the-meaning-of-following-inline-assembly-code>

System Calls: how do they work?

- Execute internal interrupt (**int instruction**)
 - syscall handler (struct intr_frame *f)
- Stack pointer: **f->esp**
- Program pointer: **f->eip**
- Return value just like functions (**f->eax**)
- Calling handlers
 - Pass args to handler
 - Return value to user process

System Calls: Implementation

- Read **syscall number at stack pointer**
- Dispatch a particular function to handle syscall
- Read (**validate!**) arguments (**above the stack pointer**)
 - Above the stack pointer
 - Validate pointers and buffers!
- **Syscall numbers defined in lib/syscall-nr.h**

Syscalls To Implement

- read
- write
- seek
- tell
- close
- create
- remove
- open
- filesize

File syscall



halt

exec

exit

wait

Process syscall



System Call: File System

- Many syscalls involve file system functionality
- Simple filesystem implement is provided: *filesys.h*, *file.h*
 - No need to modify it, but familiarize yourself
- File system is not **thread-safe**!
 - Use a coarse-grained lock to protect it
- Syscalls take **file descriptors** as args
 - Pintos represents files with struct file*
 - You must design the mapping

System Calls: Processes(1)

- **Generally, these syscalls require the most design and implementation time**
- **pid_t exec(const char *cmd line)**
 - Similar to UNIX fork() + execve()
 - Creates a child process
 - Returns after the new process has been created
 - Creation is successful if child has successfully loaded its executable and there is a thread ready to run

System Calls: Processes(2)

- **int wait (pid_t pid)**
 - parent must block until child process pid exits
 - returns exit status of the child
 - must work if child has ALREADY exited
 - must fail if it has already been called on child before
 - you may need to consider many race conditions
- **void exit (int status)**
 - exit with status and free resources
 - process termination message
 - parent must be able to retrieve status via wait

System Calls: Security

- **How does system recover from null pointer segfault in user program?**
 - kill user process, life goes on
- **What about in kernel space?**
 - Verify all user-passed memory references (pointers, buffers, strings)
 - Kill user program if passed illegal addresses

Denying Writes To Executables(Extra Credit)

- **Executables are files like any other**
- **Pintos should not allow code that is currently running to be modified**
 - Use `file_deny_write()` to prevent writes to an open file
 - Closing a file will re-enable writes
 - Keep executable open as long as the process is running

Outline

- User Programs In Pintos
- An Overview of Project 2
- Getting Start
- Tips

Getting Started

- **Lab 2 does not depend on Lab 1**
 - You can either build on your lab1 submission or start from beginning
- **Lab 3 and lab 4 are built on top of lab2**
 - Any design defects in lab 2 might affect lab3 and lab4


Getting Started: File System Setup

- You need to format a file system to store user programs
- Create a simulated disk called **filesys.dsk** with a 2MB Pintos file system partition, and then copy programs and run them
 - Make disk: pintos-mkdisk filesys.dsk --filesystem-size=2
 - Format disk: pintos -- -f -q
 - Copy program: pintos -p ../../examples/echo -a echo -- -q
 - Run program: pintos -q run 'echo x'

Getting Started: Implement this first! (1)

- **Argument passing:** change `*esp = PHYS_BASE;` to `*esp = PHYS_BASE - 12;`
 - Allows running programs with no arguments
 - Change again to correct implementation later
- **User memory access**
 - All system calls need to read user memory
- **System call infrastructure**
 - Read system call number from the user stack and dispatch to a handler

Why $*esp = PHYS_BASE - 12$?



Getting Started: Implement this first! (2)

- **Exit system call**
 - Write system call for STDOUT
- **Temporarily change process_wait to an infinite loop so pintos doesn't immediately power off**
- **Refine your implementation and pass the test**

Outline

- User Programs In Pintos
- An Overview of Project 2
- Getting Start

 **Tips**

General Tips

- **Key to implement lab2: understand the user program**
 - 80x86 Calling Convention
 - Program Startup Details
 - System Call Details
- **Read the design doc together, make sure every member in your group understand the user program**
- **Follow the suggested order of implementation!**
- **Be brave in modifying original definitions**

Debugging Tips

- **If you're confused about why a test is failing, read the source code in tests/userprog**
- **Read the system call APIs carefully, and make sure you validate all user memory addresses**

Common Errors(1)

- **My string is modified after being strtok_r()!**
 - strtok_r() modifies the string, so copy it first
 - be careful when allocating memory for copied buffer, allocating a large buffer = kernel PANIC!
- **hex_dump() prints nothing like it is supposed to be!**
 - check your user page layout and double check how it would be copied to kernel page, also did you specify the right address to print

Common Errors(2)

- **Process terminates before it prints anything!**
 - before you implement sys_wait(), use a while(1) loop to hang main thread so you can see output from user programs
- **Any program with arguments will fail!**
 - use `*esp = PHYS_BASE - 12;` for now
 - or you can implement arguments passing first (~1 hour)

Security Tips

- Cast struct file * to int, and use it as the file descriptor? Use struct thread * as pid_t?
 - info leak
- write() can be used to dump kernel memory to a file
 - Forget to check kernel memory boundary?
 - read() can be used to overwrite kernel memory
- User program takes over kernel!